

public-key cryptography

slides

bit.ly/cs161-disc

feedback

bit.ly/extended-feedback

~~hack~~ mitigation of the day

~~hack~~ mitigation of the day

- [Google introduces client-side encryption to Gmail and Calendar](#)

~~hack~~ mitigation of the day

- [Google introduces client-side encryption to Gmail and Calendar](#)
 - data encrypted on your computer before sent to Google (via HTTPS)

~~hack~~ mitigation of the day

- [Google introduces client-side encryption to Gmail and Calendar](#)
 - data encrypted on your computer before sent to Google (via HTTPS)
- “...customers have sole control over their keys

~~hack~~ mitigation of the day

- Google introduces client-side encryption to Gmail and Calendar
 - data encrypted on your computer before sent to Google (via HTTPS)
- “...customers have sole control over their keys

SERVICE	DATA THAT'S CLIENT-SIDE ENCRYPTED	DATA THAT'S NOT CLIENT-SIDE ENCRYPTED
Google Drive	<ul style="list-style-type: none">Files created with Google Docs Editors (documents, spreadsheets, presentations)Uploaded files, like PDFs and Microsoft Office files	<ul style="list-style-type: none">File titleFile metadata, such as owner, creator, and last-modified timeDrive labels (also called Drive metadata)Linked content that's outside of Docs or Drive (for example, a YouTube video linked from a Google document)User preferences, such as Docs header styles
Gmail	<ul style="list-style-type: none">Email body, including inline imagesAttached files Note: Attaching client-side encrypted Drive files isn't yet supported	<ul style="list-style-type: none">Email header, including subject, timestamps, and recipients lists
Google Calendar	<ul style="list-style-type: none">Event descriptionAttached Drive files (if CSE for Drive is turned on)Meet audio and video streams (if CSE for Meet is turned on)	<p>Any content other than the event description, attachments, and Meet data, such as:</p> <ul style="list-style-type: none">Event titleEvent starting and ending timesAttendees listBooked roomsJoin by phone numbersLink for Meet
Google Meet	<ul style="list-style-type: none">Audio streamsVideo streams (including screen sharing)	<ul style="list-style-type: none">Any data other than audio and video streams

general questions, concerns, etc.

public-key cryptography

public-key cryptography

- every individual has two keys

public-key cryptography

- every individual has two keys
 - public key: known to everybody

public-key cryptography

- every individual has two keys
 - public key: known to everybody
 - private key
 - corresponds to public key

public-key cryptography

- every individual has two keys
 - public key: known to everybody
 - private key
 - corresponds to public key
- no longer need to assume alice and bob share key

public-key cryptography

- every individual has two keys
 - public key: known to everybody
 - private key
 - corresponds to public key
- no longer need to assume alice and bob share key
- much slower than symmetric-key cryptography
 - relies on number theory calculations

el gamal

el gamal

- diffie-hellman, but with simultaneous encrypt

-

el gamal

- diffie-hellman, but with simultaneous encrypt
- Bob: private key b and public key $B = g^b \bmod p$

el gamal

- diffie-hellman, but with simultaneous encrypt
- Bob: private key b and public key $B = g^b \text{ mod } p$
- Alice: random r and computes $R = g^r \text{ mod } p$

el gamal

- diffie-hellman, but with simultaneous encrypt
- Bob: private key b and public key $B = g^b \bmod p$
- Alice: random r and computes $R = g^r \bmod p$
- Alice sends $C_1 = R$, $C_2 = M \times B^r \bmod p$

el gamal

- diffie-hellman, but with simultaneous encrypt
- Bob: private key b and public key $B = g^b \bmod p$
- Alice: random r and computes $R = g^r \bmod p$
- Alice sends $C_1 = R$, $C_2 = M \times B^r \bmod p$
- Bob: $C_2 \times C_1^{-b} = M \times B^r \times R^{-b} = M \times g^{br} \times g^{-br} = M \bmod p$

RSA - KeyGen()

RSA - KeyGen()

- random large primes p and q

RSA - KeyGen()

- random large primes p and q
- Compute $N = pq$

RSA - KeyGen()

- random large primes p and q
- Compute $N = pq$
- Choose e

RSA - KeyGen()

- random large primes p and q
- Compute $N = pq$
- Choose e
 - Requirement: e is relatively prime to $(p - 1)(q - 1)$

RSA - KeyGen()

- random large primes p and q
- Compute $N = pq$
- Choose e
 - Requirement: e is relatively prime to $(p - 1)(q - 1)$
 - Requirement: $2 < e < (p - 1)(q - 1)$

RSA - KeyGen()

- random large primes p and q
- Compute $N = pq$
- Choose e
 - Requirement: e is relatively prime to $(p - 1)(q - 1)$
 - Requirement: $2 < e < (p - 1)(q - 1)$
- Compute $d = e^{-1} \bmod (p - 1)(q - 1)$
 - Extended Euclid's algorithm

RSA - KeyGen()

- random large primes p and q
- Compute $N = pq$
- Choose e
 - Requirement: e is relatively prime to $(p - 1)(q - 1)$
 - Requirement: $2 < e < (p - 1)(q - 1)$
- Compute $d = e^{-1} \bmod (p - 1)(q - 1)$
 - Extended Euclid's algorithm
- Public key: N and e

RSA - KeyGen()

- random large primes p and q
- Compute $N = pq$
- Choose e
 - Requirement: e is relatively prime to $(p - 1)(q - 1)$
 - Requirement: $2 < e < (p - 1)(q - 1)$
- Compute $d = e^{-1} \bmod (p - 1)(q - 1)$
 - Extended Euclid's algorithm
- Public key: N and e
- Private key: d

RSA - encrypt

Alice

has access to Bob's
public keys N, e

Bob

has private key d ,
public keys N, e

RSA - encrypt

message




Alice

has access to Bob's
public keys N, e

Bob

has private key d ,
public keys N, e

RSA - encrypt



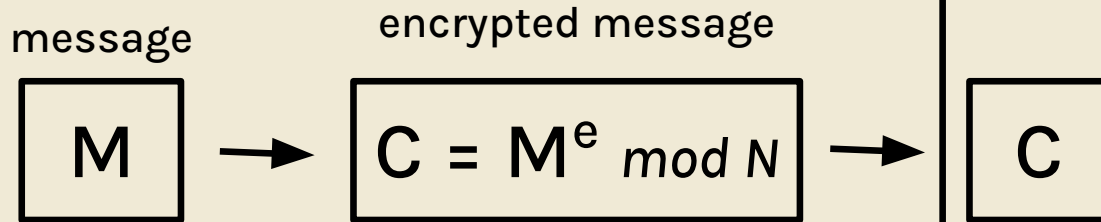
Alice

has access to Bob's
public keys N, e

Bob

has private key d ,
public keys N, e

RSA - encrypt



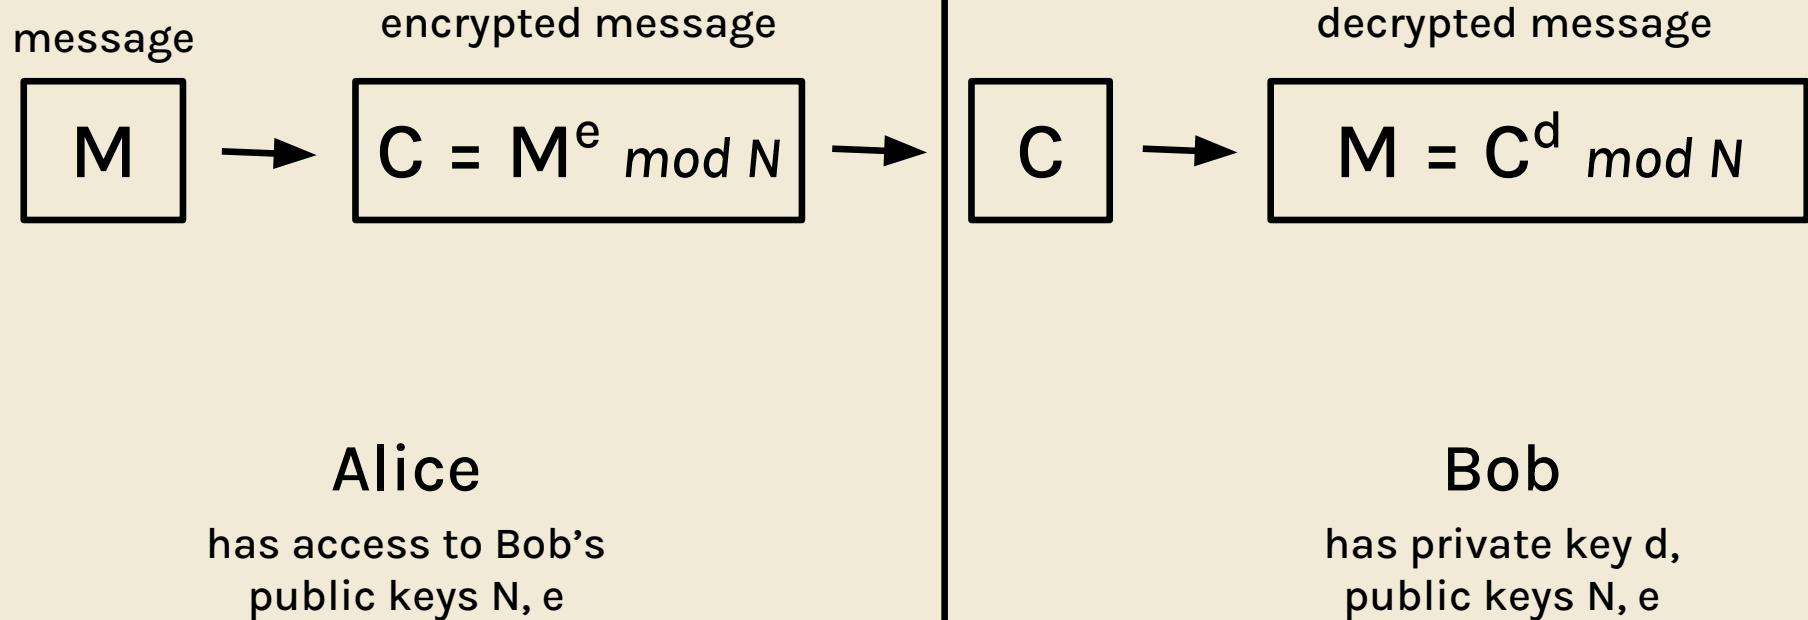
Alice

has access to Bob's
public keys N, e

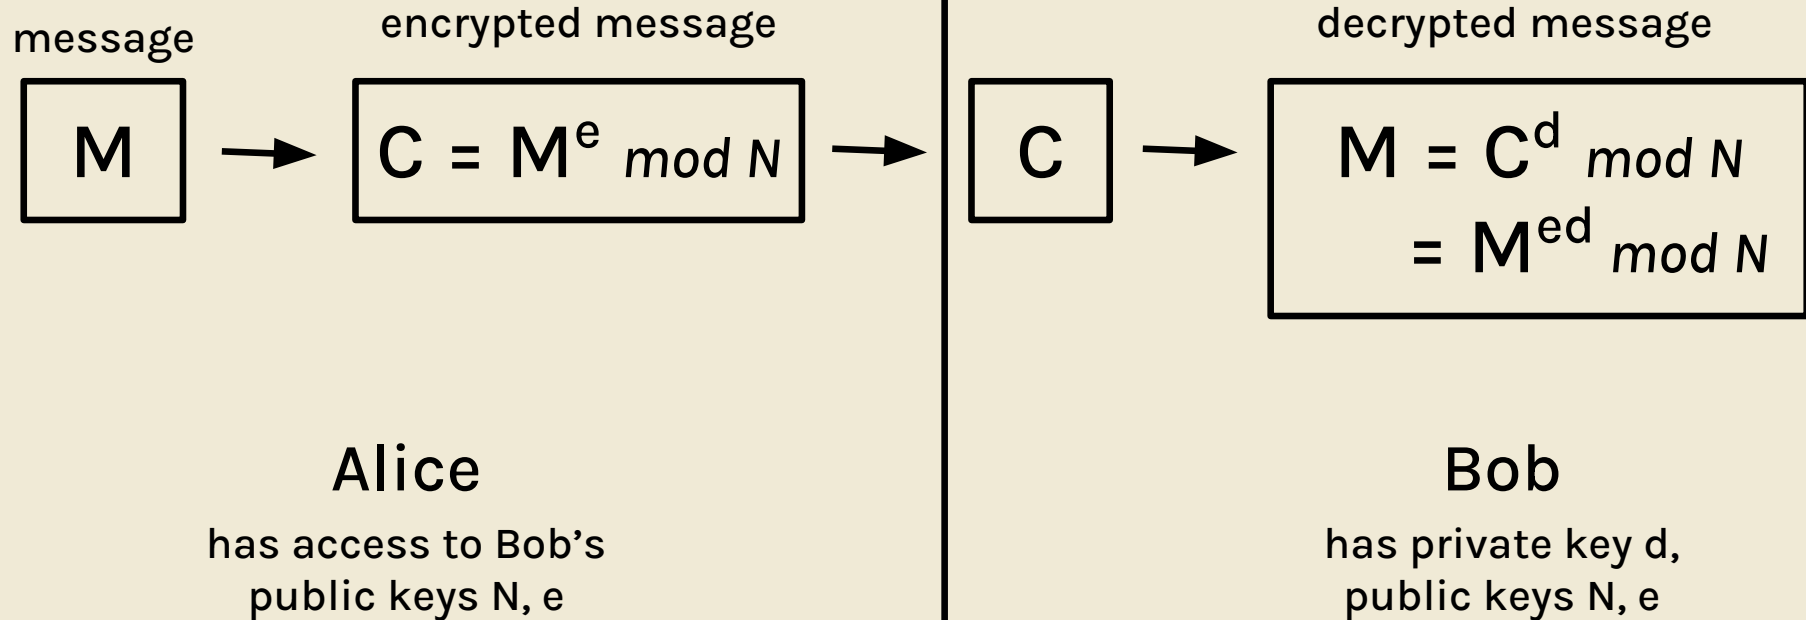
Bob

has private key d ,
public keys N, e

RSA - encrypt



RSA - encrypt



public key's downside

public key's downside

- relies on accurate public keys given

public key's downside

- relies on accurate public keys given
 - is public key cryptography secure against man-in-the-middle attacks?

public key's downside

- relies on accurate public keys given
 - is public key cryptography secure against man-in-the-middle attacks?
 - no! what if mallory gives alice her public key instead of bob's?

distributing public keys

distributing public keys

- bob can just sign his public key to ensure it's his!

distributing public keys

- bob can just sign his public key to ensure it's his!
 - problems with this?

distributing public keys

- bob can just sign his public key to ensure it's his!
 - problems with this?
 - must have bob's PK to verify signature

distributing public keys

- bob can just sign his public key to ensure it's his!
 - problems with this?
 - must have bob's PK to verify signature
- can't gain trust without trusting anything

distributing public keys

- bob can just sign his public key to ensure it's his!
 - problems with this?
 - must have bob's PK to verify signature
- can't gain trust without trusting anything
- **trust anchor**: a root of trust—implicit trust

trust-on-first-use

trust-on-first-use

trust-on-first-use

- trust the user's PK on first communication

trust-on-first-use

- trust the user's PK on first communication
- warn user if PK changes

trust-on-first-use

- trust the user's PK on first communication
- warn user if PK changes
- idea: attacks are infrequent, probably won't happen the first time

trust-on-first-use

- trust the user's PK on first communication
- warn user if PK changes
- idea: attacks are infrequent, probably won't happen the first time
- used by SSH

certificates

certificates

- signed endorsement of someone's public key

certificates

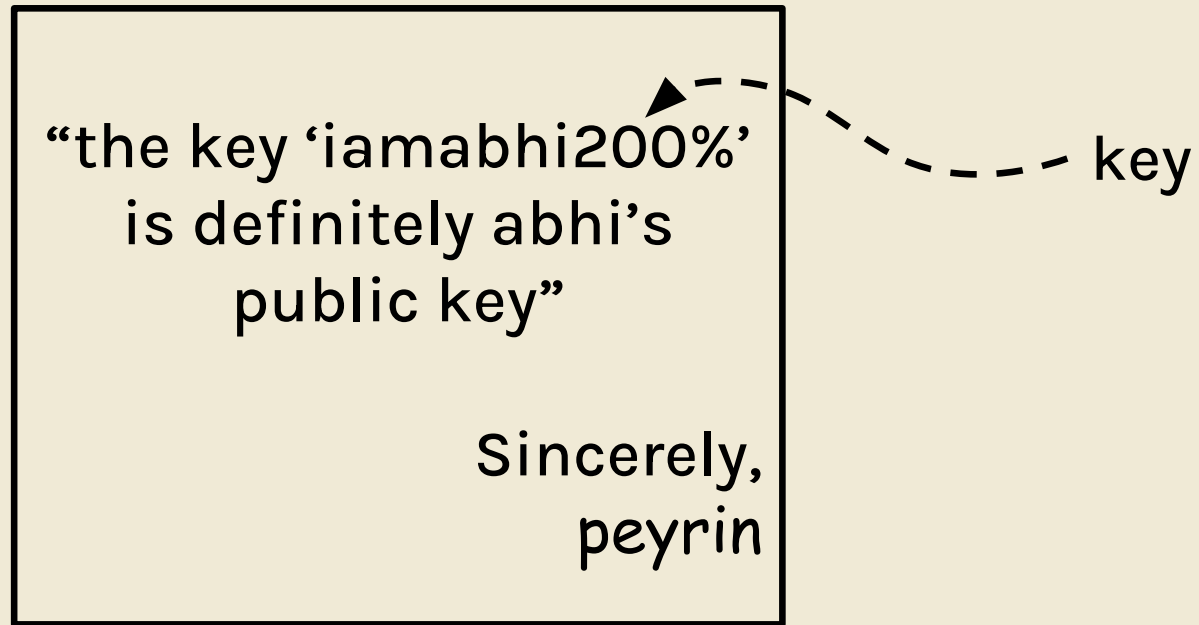
- signed endorsement of someone's public key

“the key ‘iamabhi200%’
is definitely abhi’s
public key”

Sincerely,
peyrin


certificates

- signed endorsement of someone's public key




certificates

- signed endorsement of someone's public key



certificates

- signed endorsement of someone's public key



certificates

- signed endorsement of someone's public key
- encryption under a public key PK : $\{\text{"Message"}\}_{PK}$

certificates

- signed endorsement of someone's public key
- encryption under a public key $PK: \{\text{"Message"}\}_{PK}$
- signing with a private key $SK: \{\text{"Message"}\}_{SK^{-1}}$

certificates

- signed endorsement of someone's public key
- encryption under a public key PK : $\{\text{"Message"}\}_{PK}$
- signing with a private key SK : $\{\text{"Message"}\}_{SK^{-1}}$
 - signed message = message & signature on message (not just a signature)

certificates

- signed endorsement of someone's public key
- encryption under a public key PK : $\{\text{"Message"}\}_{PK}$
- signing with a private key SK : $\{\text{"Message"}\}_{SK^{-1}}$
 - signed message = message & signature on message (not just a signature)
- if we trust EvanBot and he sends $\{\text{"Bob's public key is } PK_B\}\}_{SK_E^{-1}}$, we trust this certificate

models of certificate trust

models of certificate trust

- trusted directory

models of certificate trust

- trusted directory
 - central directory with *PKTD*, *SKTD*

models of certificate trust

- trusted directory
 - central directory with *PKTD*, *SKTD*
 - we ask the TD for bob's PK

models of certificate trust

- trusted directory
 - central directory with PK_{TD} , SK_{TD}
 - we ask the TD for bob's PK
 - TD says {"Bob's public key is PK_B "} SK_{TD}^{-1}

models of certificate trust

- trusted directory
 - central directory with PK_{TD} , SK_{TD}
 - we ask the TD for bob's PK
 - TD says {"Bob's public key is PK_B "} SK_{TD}^{-1}
 - trust assumptions

models of certificate trust

- trusted directory
 - central directory with PK_{TD} , SK_{TD}
 - we ask the TD for bob's PK
 - TD says {"Bob's public key is PK_B "} SK_{TD}^{-1}
- trust assumptions
 - we receive the right PK_{TD}

models of certificate trust

- trusted directory
 - central directory with PK_{TD} , SK_{TD}
 - we ask the TD for bob's PK
 - TD says {"Bob's public key is PK_B "} SK_{TD}^{-1}
 - trust assumptions
 - we receive the right PK_{TD}
 - the TD won't sign keys unless the owner is verified

models of certificate trust

- trusted directory
 - central directory with *PKTD*, *SKTD*
 - problems:

models of certificate trust

- trusted directory
 - central directory with *PKTD*, *SKTD*
 - problems:
 - scalability: one directory for whole world

models of certificate trust

- trusted directory
 - central directory with *PKTD*, *SKTD*
 - problems:
 - scalability: one directory for whole world
 - single point of failure: if directory is compromised, can't trust anyone


models of certificate trust

- trusted directory
 - central directory with *PKTD*, *SKTD*
- certificate authority

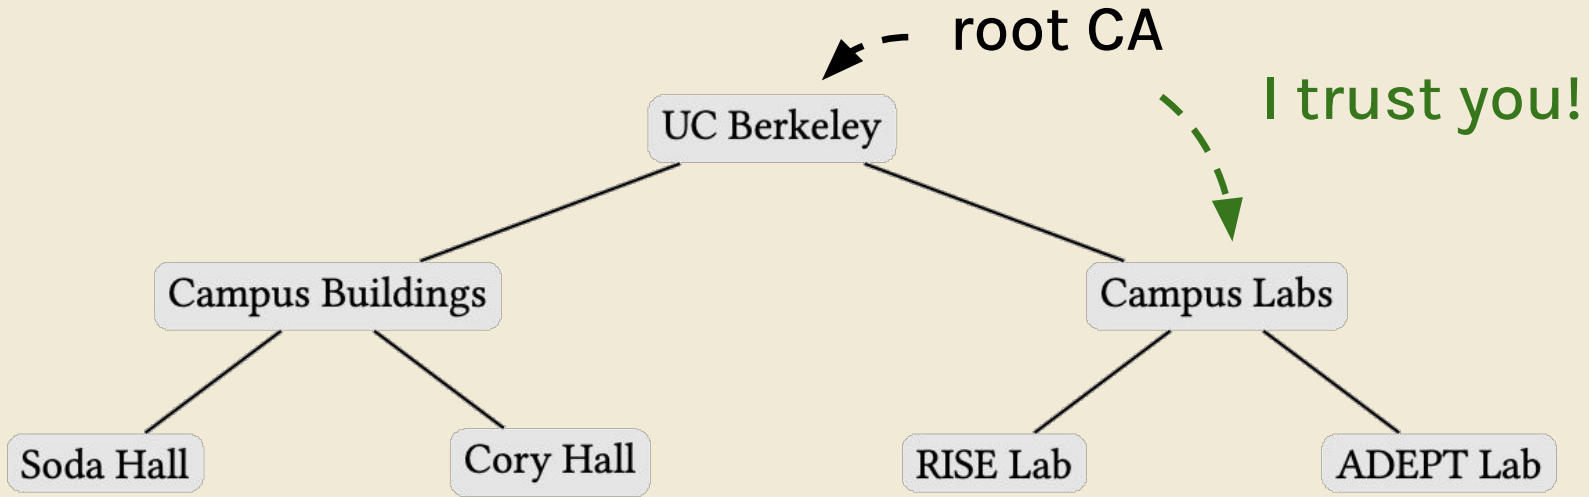
models of certificate trust

- **trusted directory**
 - central directory with *PKTD*, *SKTD*
- **certificate authority**
 - hierarchical trust: a root CA signs other CAs, and they can certificates as well


certificate authorities



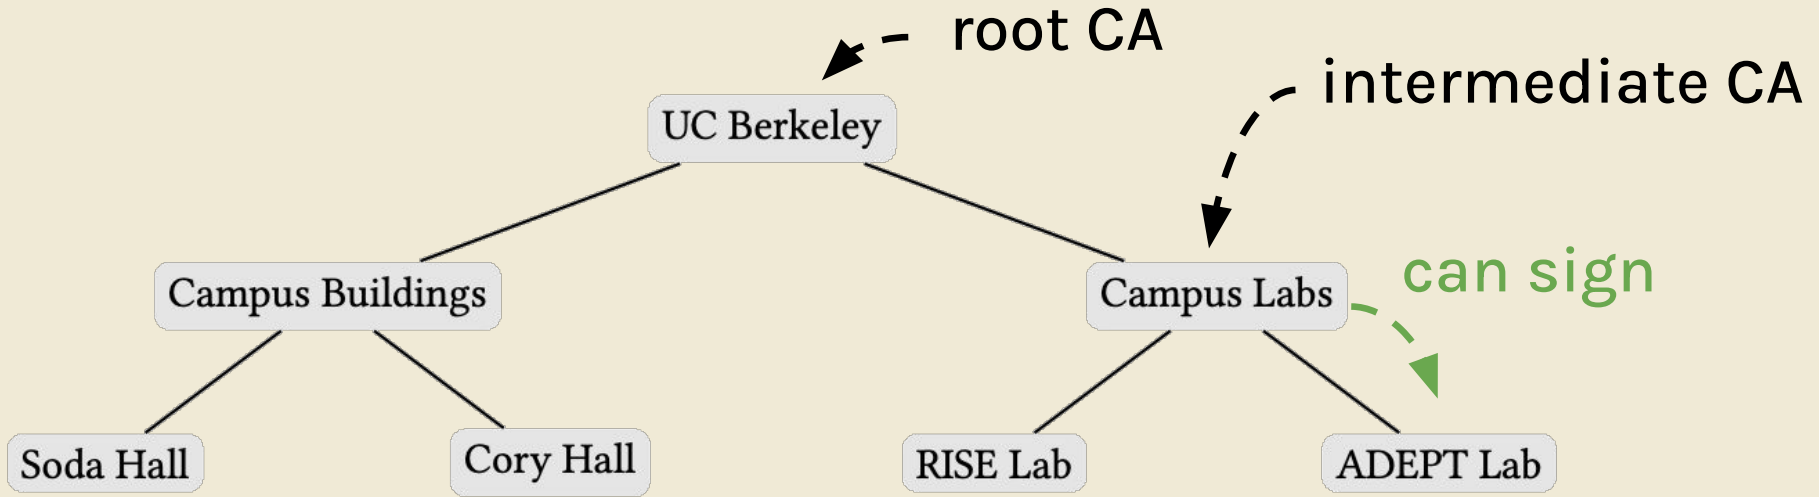
certificate authorities




certificate authorities



certificate authorities



certificate authorities



certificate revocation

certificate revocation

- when CA issues bad certificate

certificate revocation

- when CA issues bad certificate
- 1) expiration dates: when certificate expires,
re-request

certificate revocation

- when CA issues bad certificate
- 1) expiration dates: when certificate expires,
re-request
 - tradeoff?

certificate revocation

- when CA issues bad certificate
- 1) expiration dates: when certificate expires,
re-request
 - tradeoff?
- 2) certificate revocation list

certificate revocation

- when CA issues bad certificate
- 1) expiration dates: when certificate expires,
re-request
 - tradeoff?
- 2) certificate revocation list
 - {“The certificate with serial number
0xdeadbeef is now revoked”} SK_{CA}^{-1}

certificate revocation

- when CA issues bad certificate
- 1) expiration dates: when certificate expires, re-request
 - tradeoff?
 - 2) certificate revocation list
 - {“The certificate with serial number 0xdeadbeef is now revoked”} $_{SK_{CA}^{-1}}$
 - tradeoff?

hashing—a review

- $H(M)$: M is an arbitrary length message
 - output: fixed length n -bit hash
 - $\{0, 1\}^* \rightarrow \{0, 1\}^n$
- “look” random
- fast
- **one-way**: hard to find x given a y such that $H(x) = y$
- **collision-resistant**: hard to find $x \neq x'$ s.t. $H(x) = H(x')$

hashing passwords

hashing passwords

- given a password, how do I check it's the right one

hashing passwords

- given a password, how do I check it's the right one
- store the passwords in plaintext?

hashing passwords

- given a password, how do I check it's the right one
- store the passwords in plaintext?
 - an attacker can find them

hashing passwords

- given a password, how do I check it's the right one
- store the passwords in plaintext?
 - an attacker can find them
- idea: *hash* the password and store it, then check if input's hash matches stored hash

hashing passwords

- given a password, how do I check it's the right one
- store the passwords in plaintext?
 - an attacker can find them
- idea: *hash* the password and store it, then check if input's hash matches stored hash
- problem: brute-forcing passwords

hashing passwords: attacks

hashing passwords: attacks

- attacker can precompute common passwords,
e.g., $H(\text{"password"})$, $H(\text{"password123"})$, $H(\text{"peyrin"})$

hashing passwords: attacks

- attacker can precompute common passwords, e.g., $H(\text{"password"})$, $H(\text{"password123"})$, $H(\text{"peyrin"})$
 - dictionary attack: hash an entire dictionary of common passwords

hashing passwords: attacks

- attacker can precompute common passwords, e.g., $H(\text{"password"})$, $H(\text{"password123"})$, $H(\text{"peyrin"})$
 - dictionary attack: hash an entire dictionary of common passwords
- rainbow tables: an algorithm to make brute-forcing easier

brute-forcing passwords

how hard is it to brute force a 10-bit password?

password = "1011010100"

hash = $H(\text{"1011010100"}) = 158912$

brute-forcing passwords

how hard is it to brute force a 10-bit password?

password = "1011010100"

hash = $H(\text{"1011010100"}) = 158912$

only guessing 1 or 0 for 10 bits, 2 options

for each of 10 bits, $2 * 2 * 2 \dots * 2 = 2^{10}$

guesses and then hashing each guess

mitigating brute force: salts

mitigating brute force: salts

- salt: a public, random bitstring for each user

mitigating brute force: salts

- **salt:** a public, random bitstring for each user
 - should be long and random

mitigating brute force: salts

- **salt:** a public, random bitstring for each user
 - should be long and random
 - public!!! like IVs or nonces

mitigating brute force: salts

- salt: a public, random bitstring for each user
 - should be long and random
 - public!!! like IVs or nonces
- store (user, salt, $H(\text{password} || \text{salt})$) in database

mitigating brute force: salts

- **salt:** a public, random bitstring for each user
 - should be long and random
 - public!!! like IVs or nonces
- store (user, salt, $H(\text{password} || \text{salt})$) in database
- if there are M possible passwords and N users

mitigating brute force: salts

- **salt:** a public, random bitstring for each user
 - should be long and random
 - public!!! like IVs or nonces
- store (user, salt, $H(\text{password} || \text{salt})$) in database
- if there are M possible passwords and N users
 - brute force runtimes:

mitigating brute force: salts

- **salt:** a public, random bitstring for each user
 - should be long and random
 - public!!! like IVs or nonces
- store (user, salt, $H(\text{password} || \text{salt})$) in database
- if there are M possible passwords and N users
 - brute force runtimes:
 - unsalted: $O(M + N)$

mitigating brute force: salts

- salt: a public, random bitstring for each user
 - should be long and random
 - public!!! like IVs or nonces
- store (user, salt, $H(\text{password} || \text{salt})$) in database
- if there are M possible passwords and N users
 - brute force runtimes:
 - unsalted: $O(M + N)$
 - salted: $O(MN)$

mitigating brute force: slow hash

mitigating brute force: slow hash

- cryptographic hashes are fast

mitigating brute force: slow hash

- cryptographic hashes are fast
- password hashes are designed to be slow

mitigating brute force: slow hash

- cryptographic hashes are fast
- password hashes are designed to be slow
 - why?

mitigating brute force: slow hash

- cryptographic hashes are fast
- password hashes are designed to be slow
 - why?
 - guessing the wrong password should waste time

mitigating brute force: slow hash

- cryptographic hashes are fast
- password hashes are designed to be slow
 - why?
 - guessing the wrong password should waste time
 - users can't tell the difference between 0.001 and 0.1 second hashes, attackers computing thousands of hashes can

online v.s. offline attacks

online v.s. offline attacks

- online attack

online v.s. offline attacks

- online attack
 - attacker guesses password via website input

online v.s. offline attacks

- online attack
 - attacker guesses password via website input
 - defense: rate limit (certain # of guesses/min)

online v.s. offline attacks

- online attack
 - attacker guesses password via website input
 - defense: rate limit (certain # of guesses/min)
- offline attack

online v.s. offline attacks

- online attack
 - attacker guesses password via website input
 - defense: rate limit (certain # of guesses/min)
- offline attack
 - computation done by attacker, can parallelize

online v.s. offline attacks

- online attack
 - attacker guesses password via website input
 - defense: rate limit (certain # of guesses/min)
- offline attack
 - computation done by attacker, can parallelize
 - defense: salted passwords, slow hashes, strong passwords

worksheet
(on 161 website)



feedback

bit.ly/extended-feedback

slides: bit.ly/cs161-disc