

# DNS, DNSSEC

dog network service

slides

[bit.ly/cs161-disc](https://bit.ly/cs161-disc)

feedback

[bit.ly/abhifedback](https://bit.ly/abhifedback)

hack of the day

# hack of the day

- [DNS resolver insecurities found on thousands of websites out of 7000 sampled](#)

# hack of the day

- [DNS resolver insecurities found on thousands of websites out of 7000 sampled](#)
  - small/big businesses, governmental services

# hack of the day

- [DNS resolver insecurities found on thousands of websites out of 7000 sampled](#)
  - small/big businesses, governmental services
  - “...at least 25 were using static source ports... thousands of more domains using static source ports were discovered.”

# hack of the day

- [DNS resolver insecurities found on thousands of websites out of 7000 sampled](#)
  - small/big businesses, governmental services
  - “...at least 25 were using static source ports... thousands of more domains using static source ports were discovered.”
  - none of the 25 used DNSSEC, etc.

# hack of the day

- [DNS resolver insecurities found on thousands of websites out of 7000 sampled](#)
  - small/big businesses, governmental services
  - “...at least 25 were using static source ports... thousands of more domains using static source ports were discovered.”
  - none of the 25 used DNSSEC, etc.
  - cache poisoning used to hijack WordPress

general questions, concerns, etc.



# domain name system (DNS)

# domain name system (DNS)

- remember: computers use IP addresses to communicate

# domain name system (DNS)

- remember: computers use IP addresses to communicate
- but humans use domain names (google.com)

# domain name system (DNS)

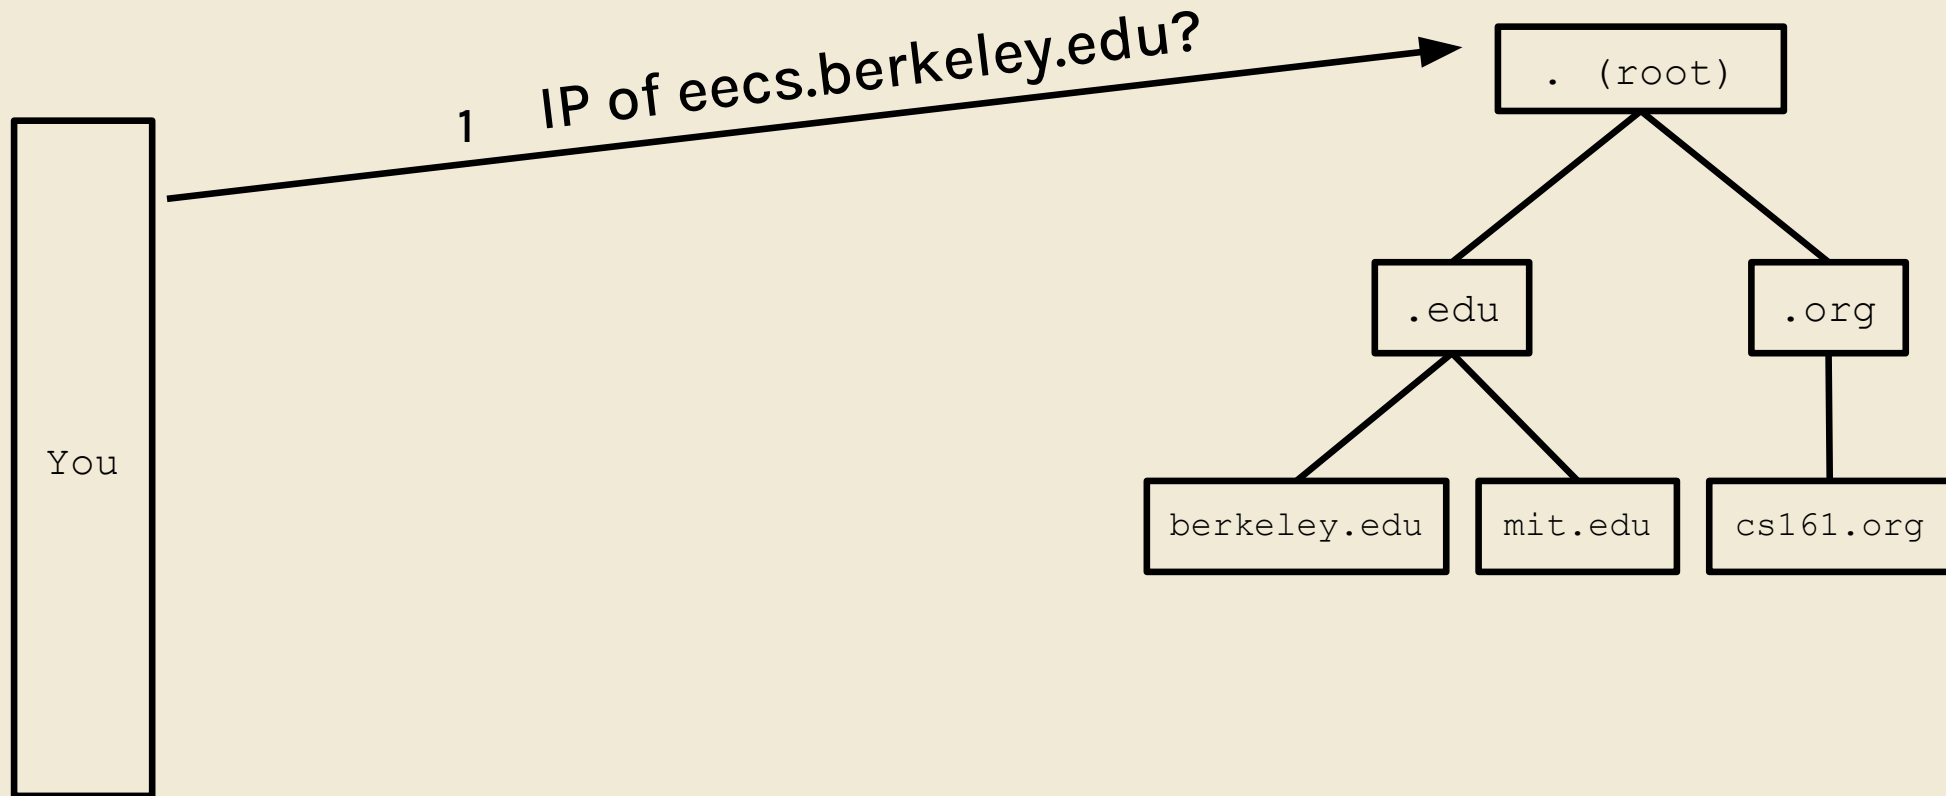
- remember: computers use IP addresses to communicate
- but humans use domain names (google.com)
- domain name -> IP address?

`www.google.com`  $\xrightarrow{\text{DNS}}$  `74.125.25.99`

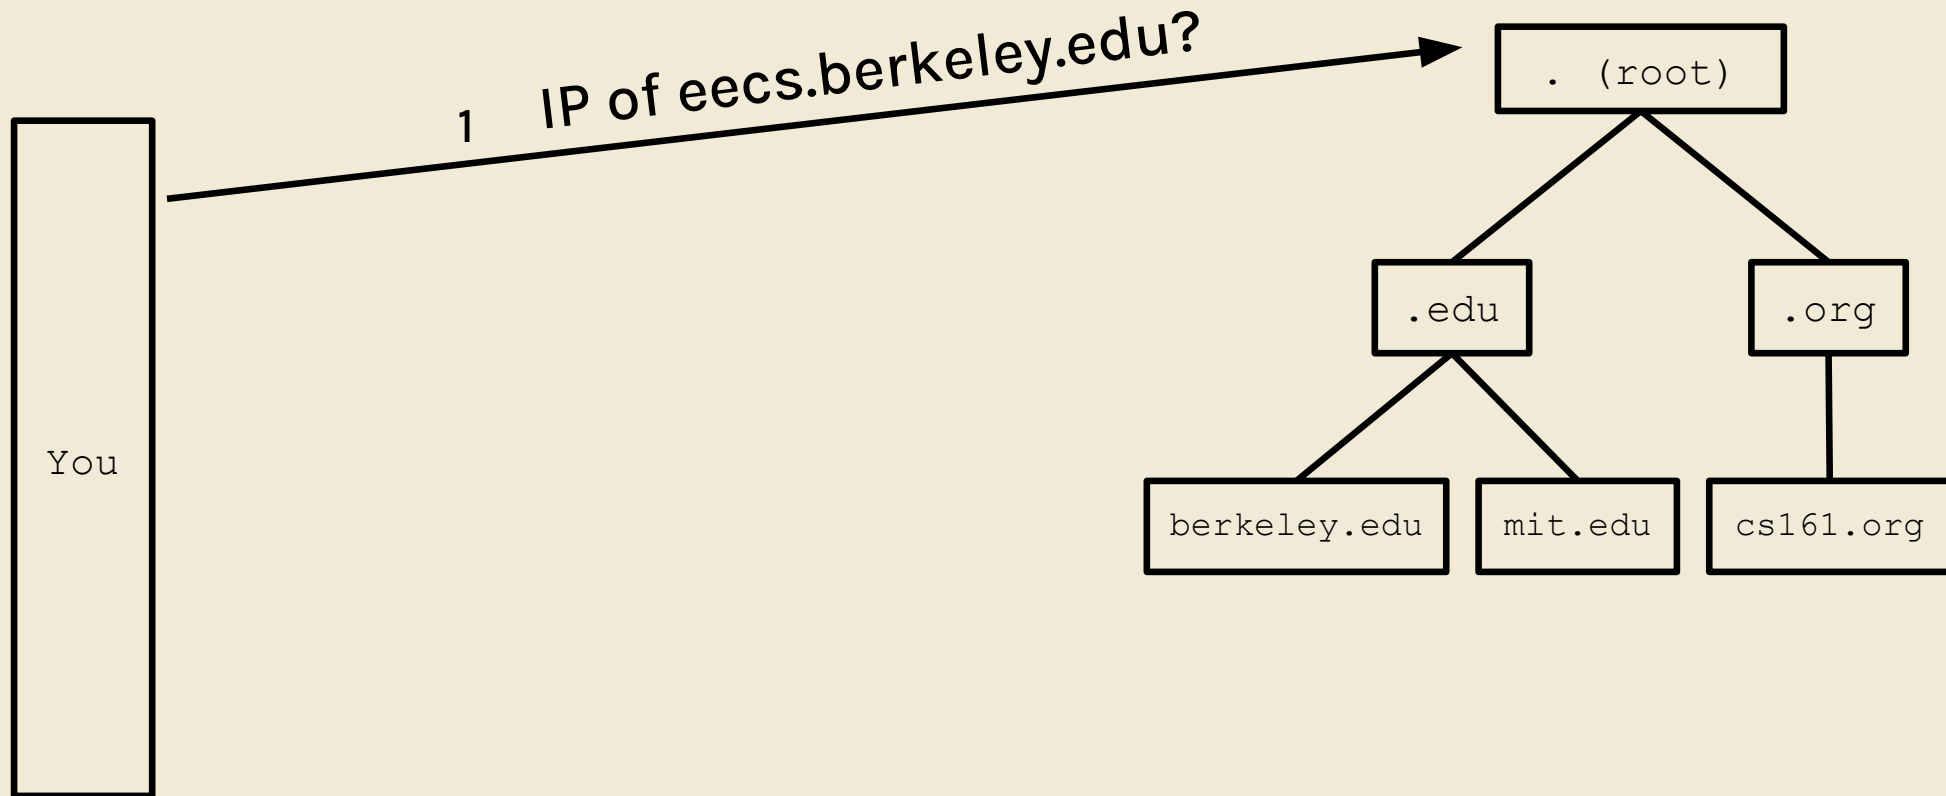
# name servers

- a server on the Internet responsible for answering DNS queries
  - “what is the IP of google.com?”
- use a hierarchy of nameservers to find the IP

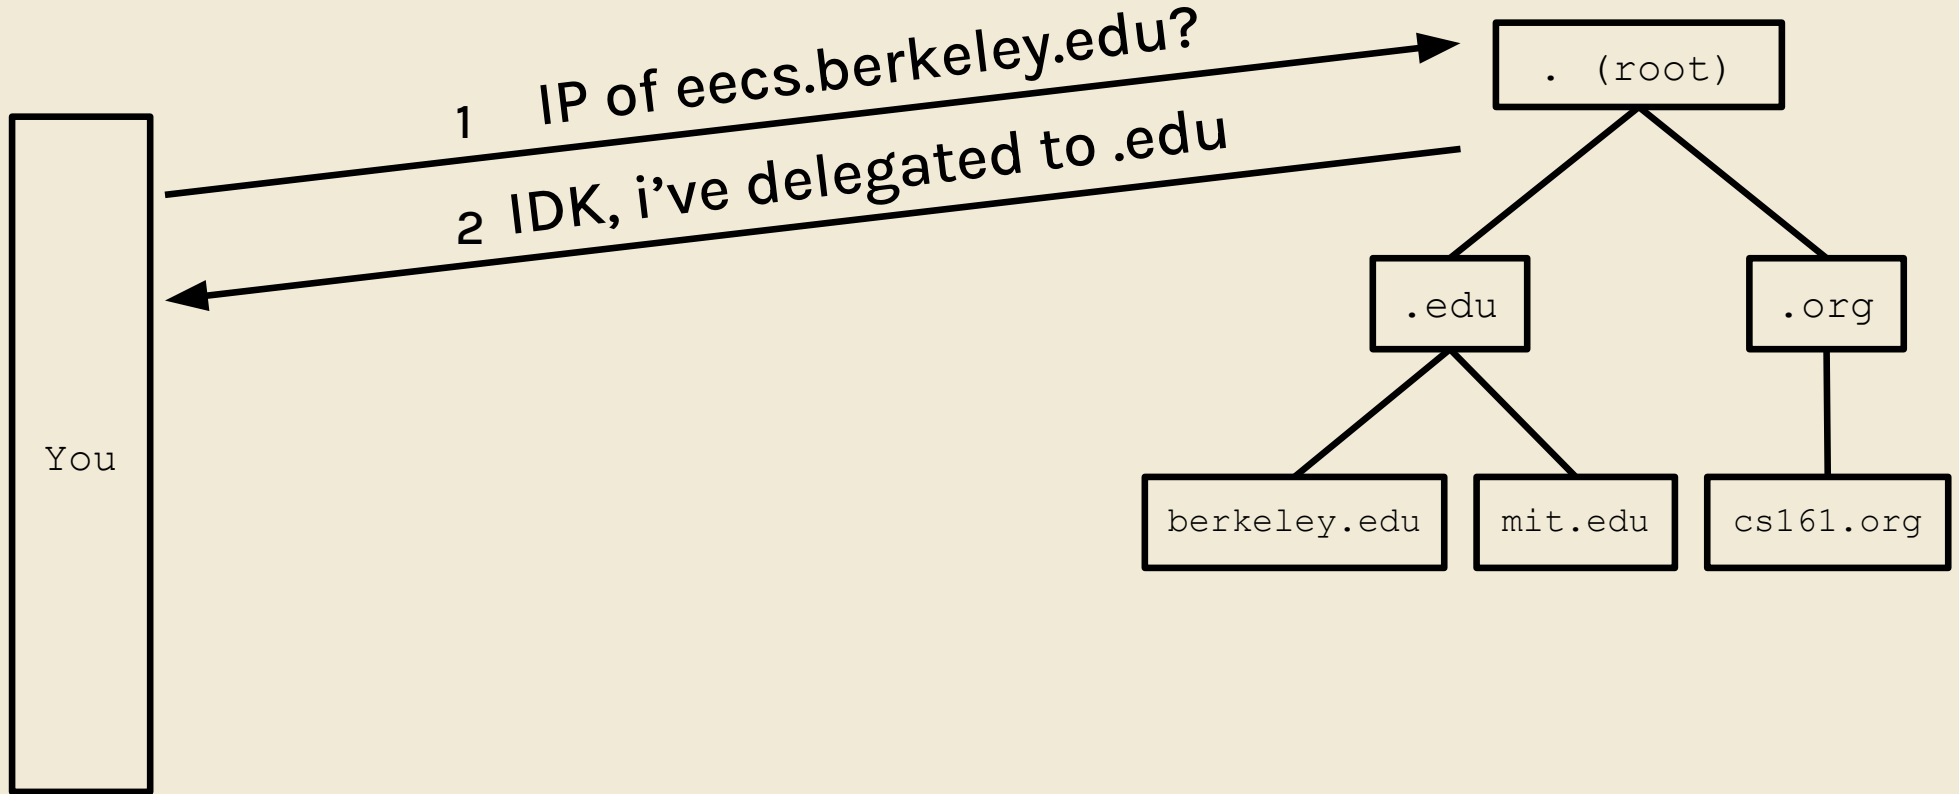
# name servers



# name servers

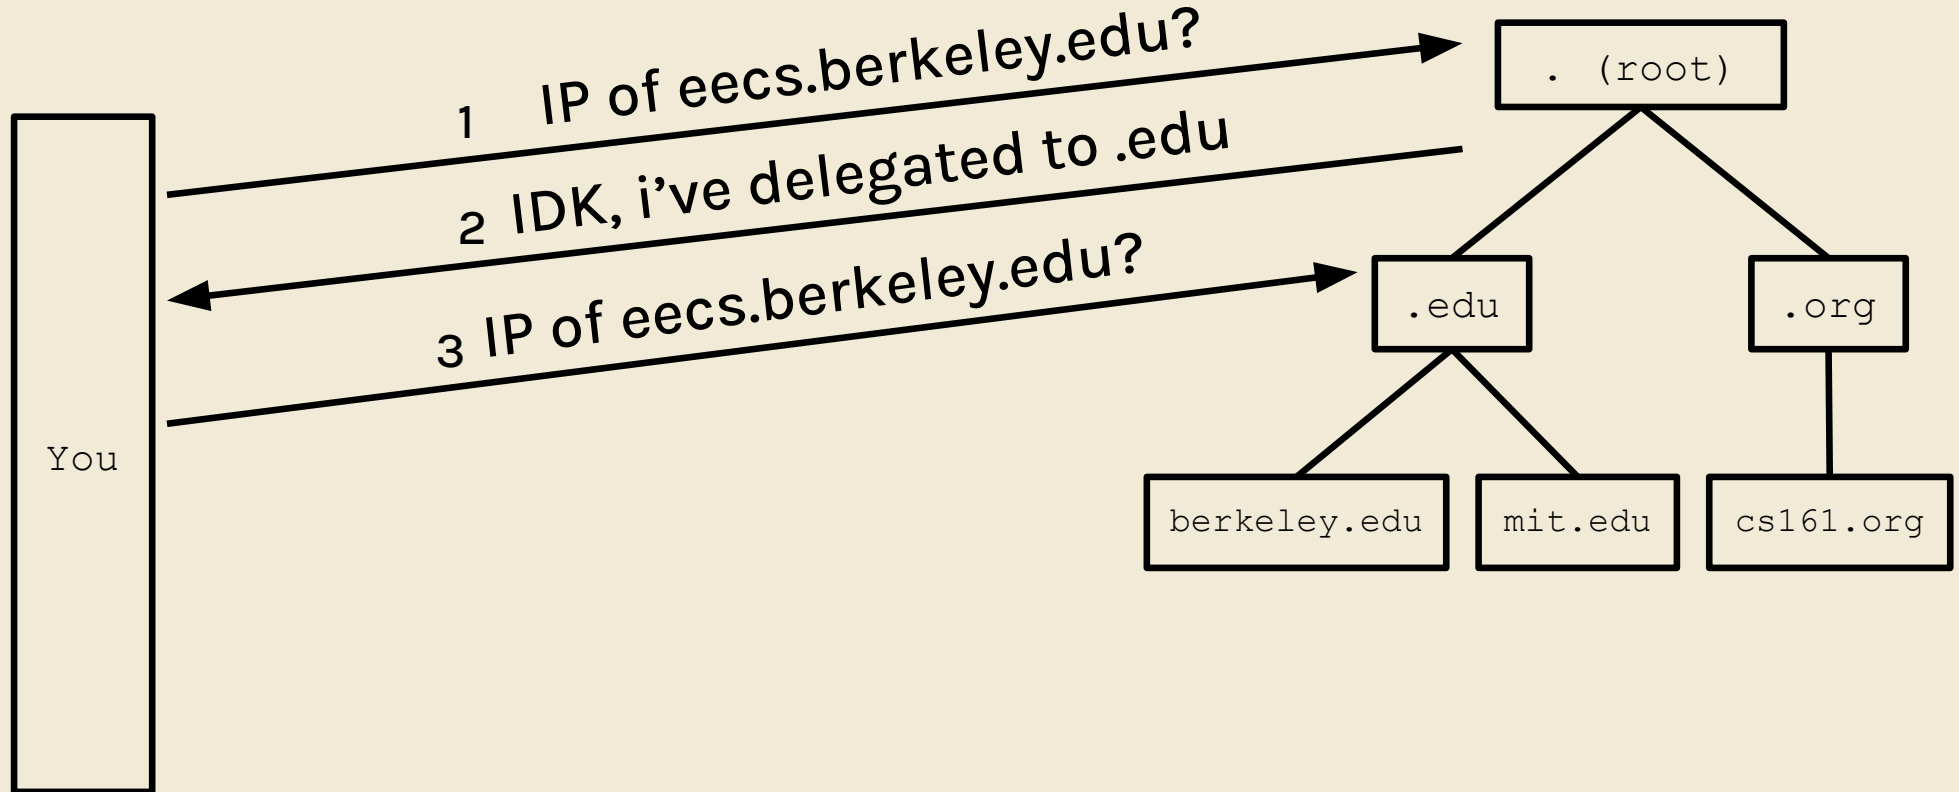


# name servers

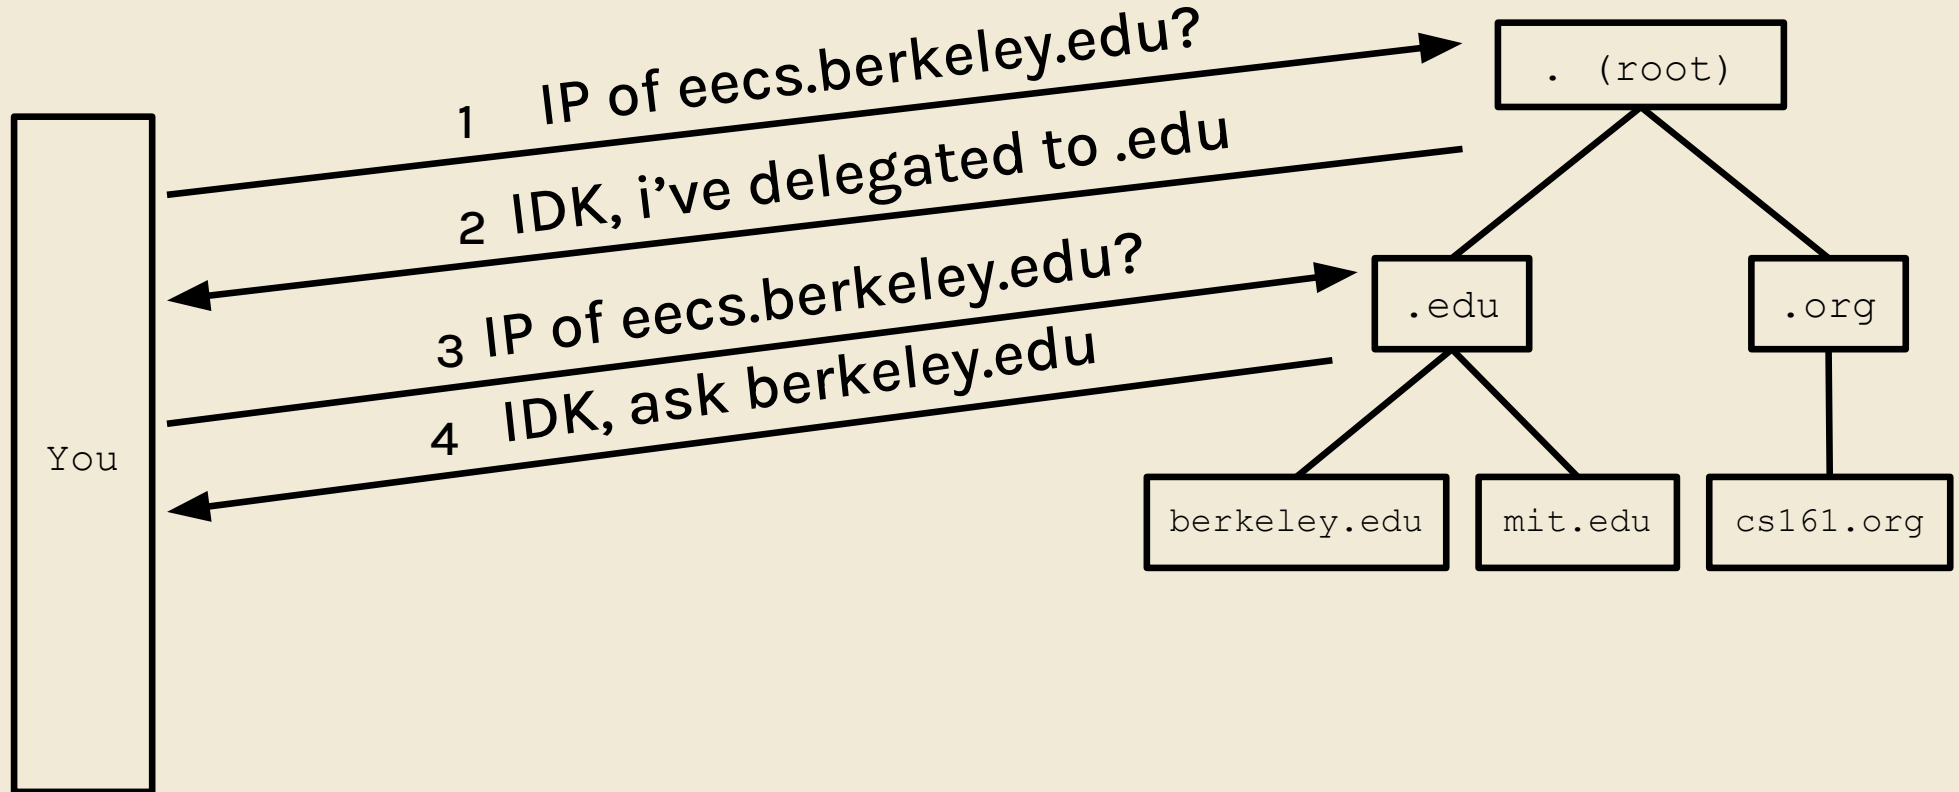




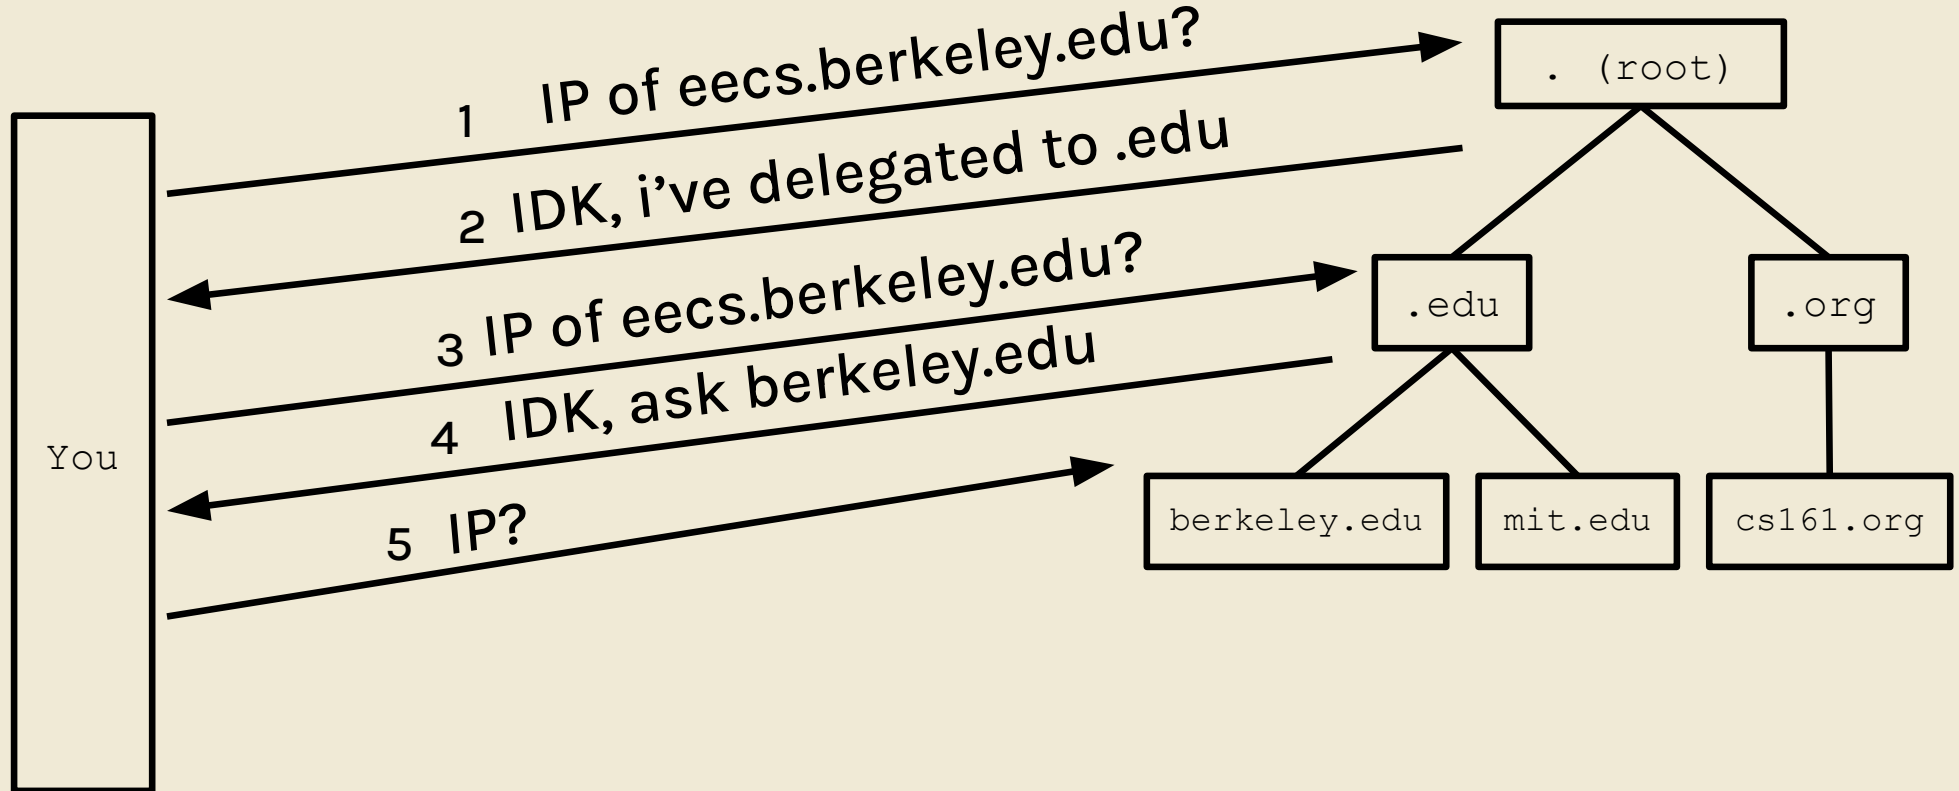
# name servers



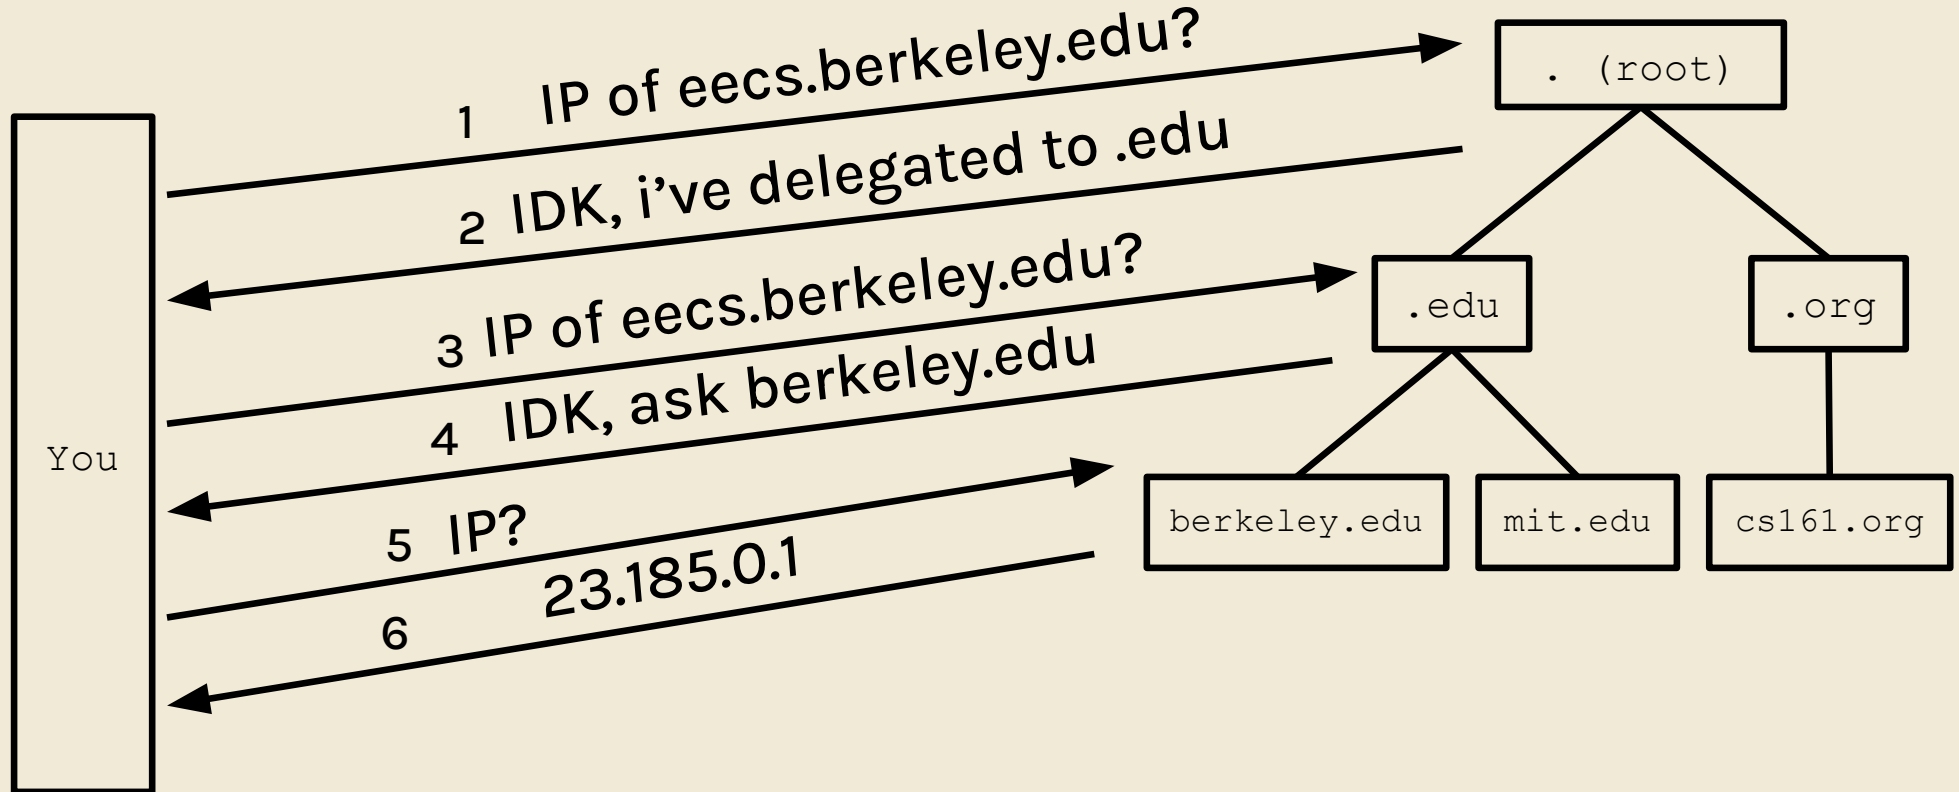
# name servers



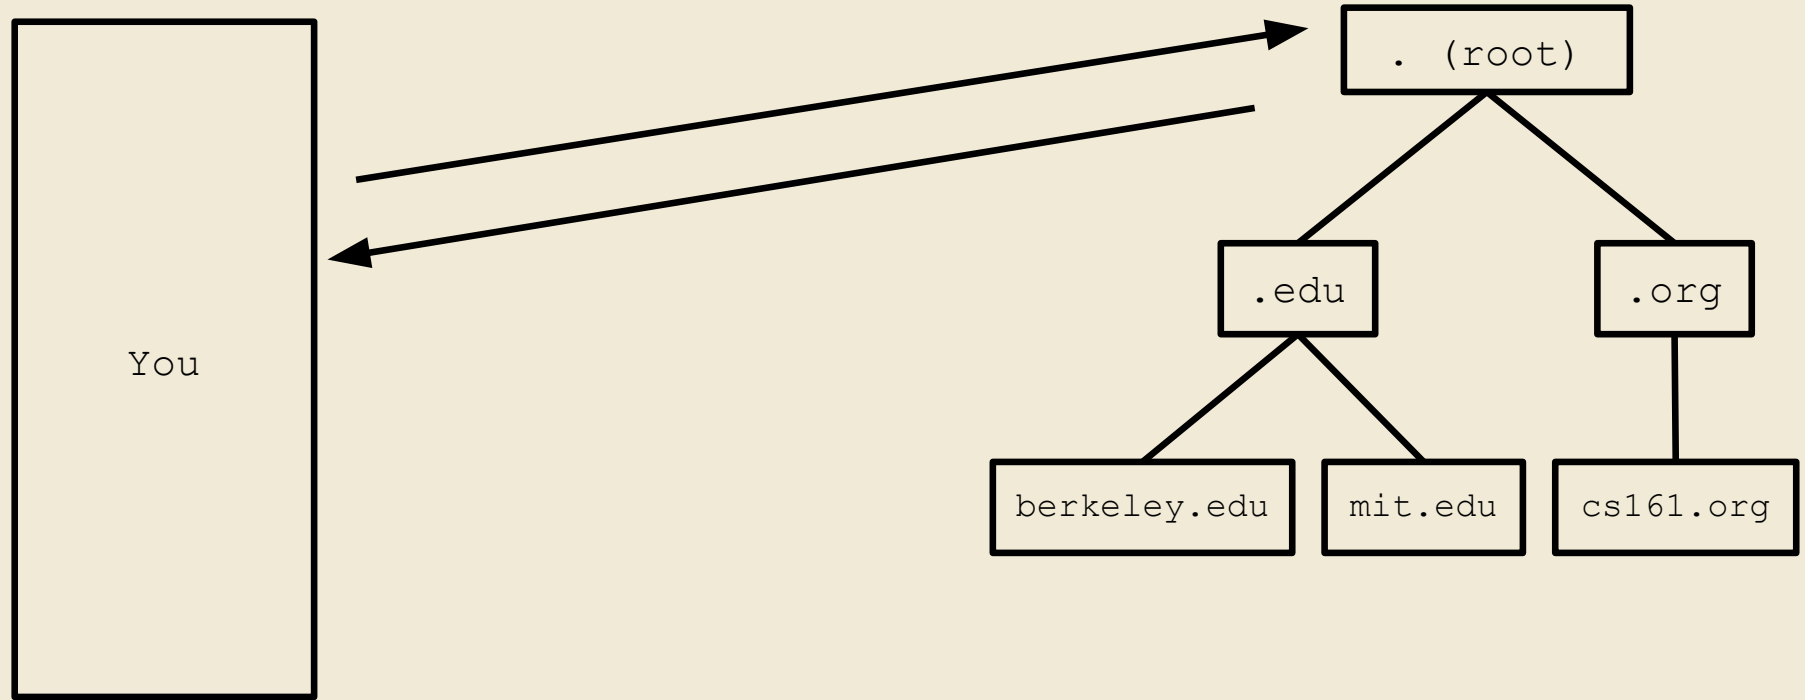
# name servers



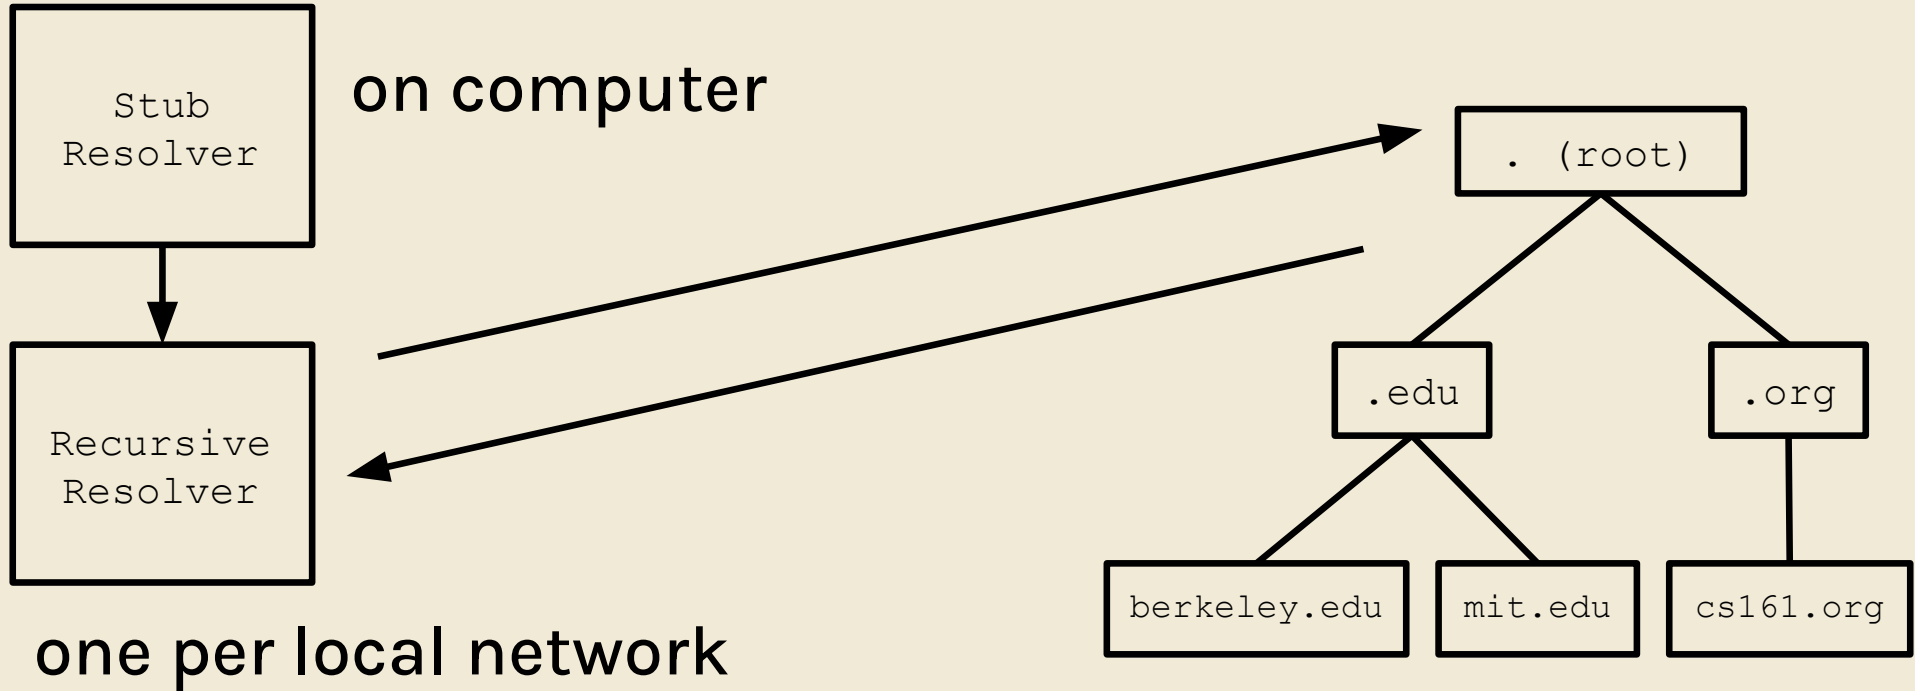
# name servers



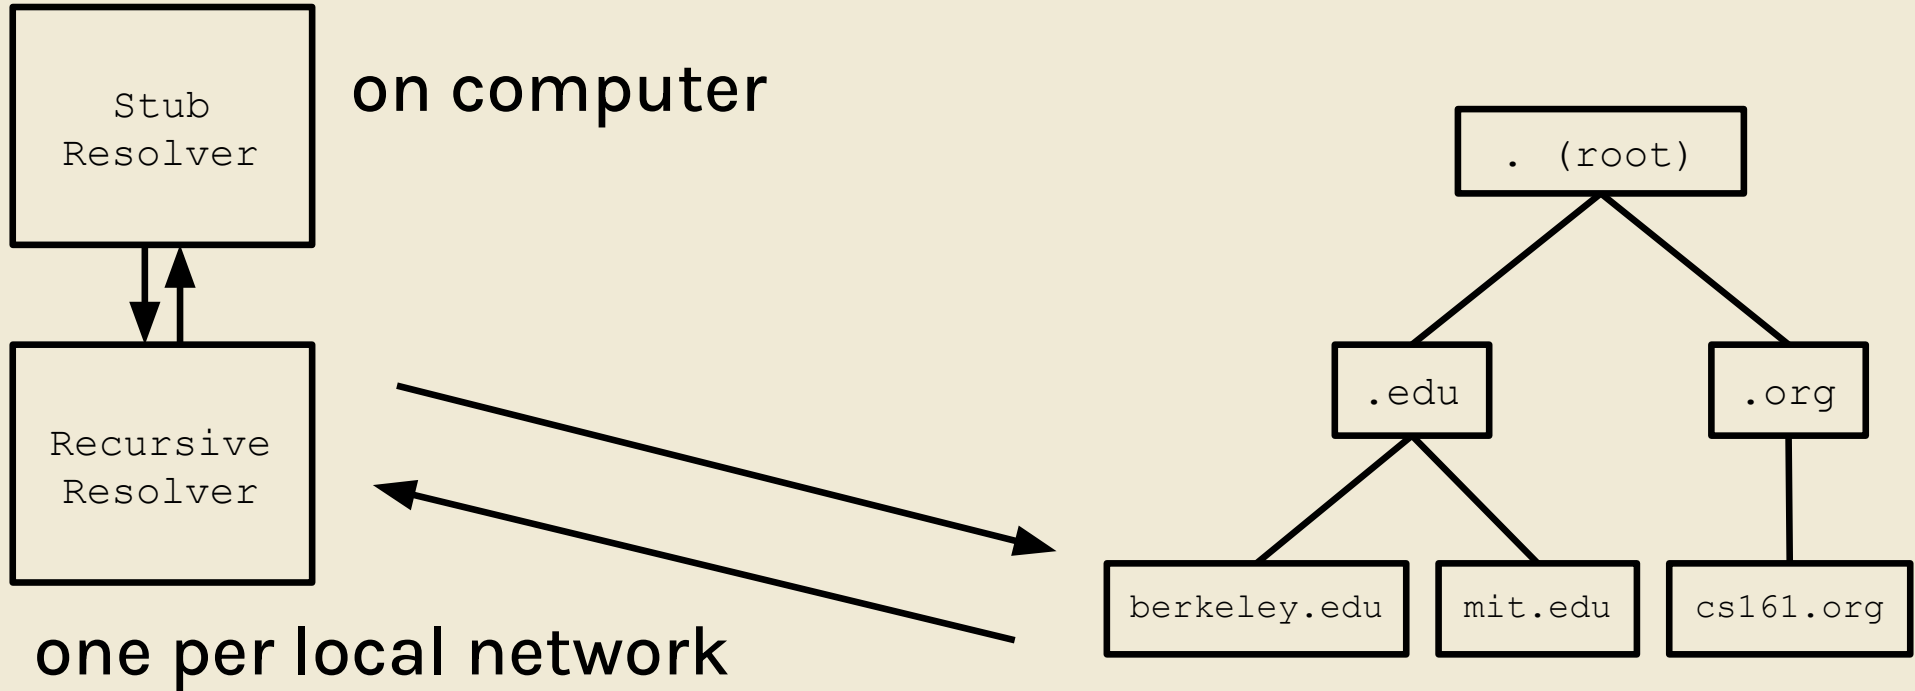
# stub/recurisve resolvers



# stub/recursive resolvers



# stub/recursive resolvers



# DNS packets

16-bit random

check for corruption

16 bits, usually 53

The diagram illustrates the structure of a DNS packet, which is encapsulated in a UDP header. The packet is divided into a 'Header' section and a 'Payload' section. The 'Header' section contains the following fields:

UDP Header	
Source Port	Destination Port
Checksum	Length
ID number	Flags
Question count	Answer count
Authority count	Additional count

The 'Payload' section contains the following records:

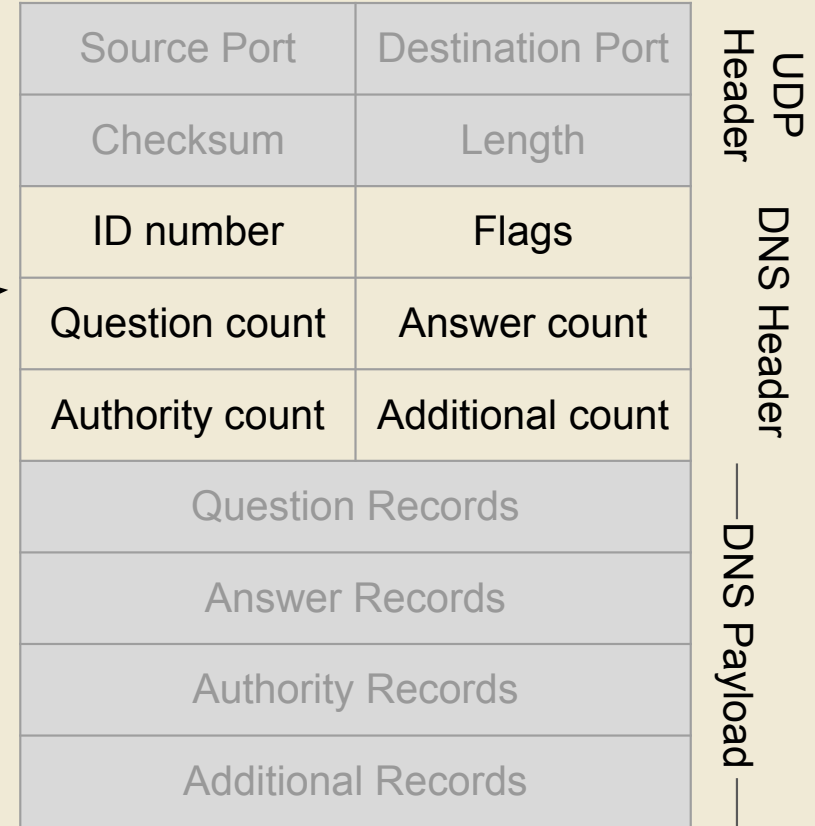
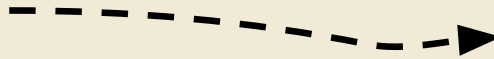
UDP Payload	
Question Records	
Answer Records	
Authority Records	
Additional Records	

Dashed arrows indicate the bit lengths for specific fields: '16-bit random' points to the 'Checksum' field; 'check for corruption' points to the 'Length' field; and '16 bits, usually 53' points to the 'ID number' field.



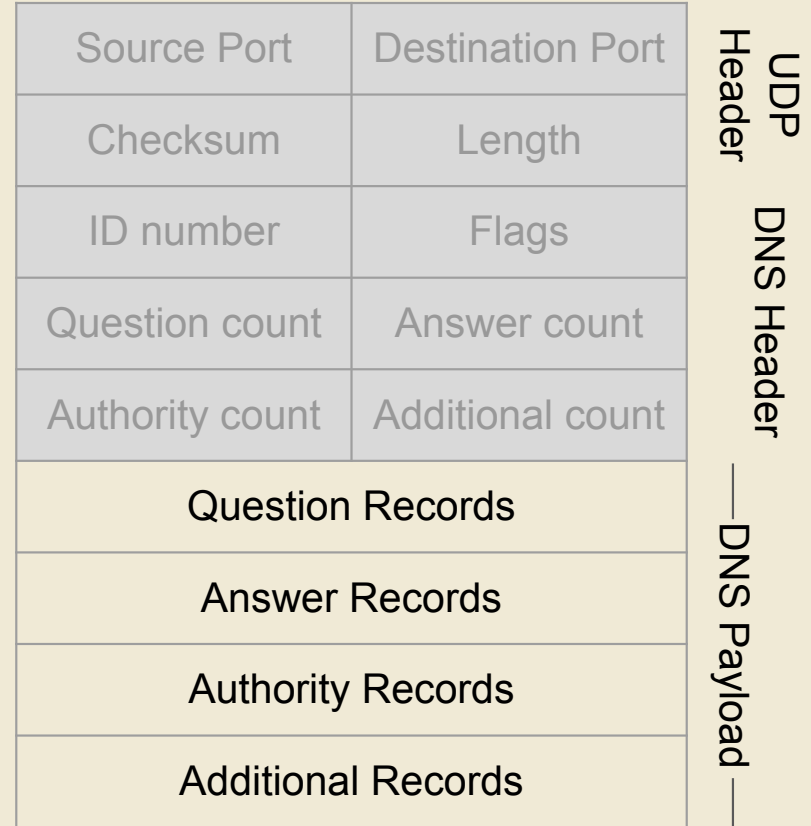
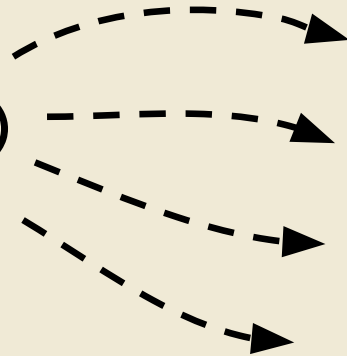
# DNS packets

query ID same as  
response ID



# DNS packets

variable number of  
resource records (RRs)



# DNS lookup walkthrough

```
$ dig +norecurse eecs.berkeley.edu @198.41.0.4
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 26114
```

```
;; flags: qr; QUERY: 1, ANSWER: 0, AUTHORITY: 13, ADDITIONAL: 27
```

```
;; QUESTION SECTION:
```

```
;eecs.berkeley.edu.          IN      A
```

```
;; AUTHORITY SECTION:
```

```
edu.          172800    IN      NS      a.edu-servers.net.
```

```
edu.          172800    IN      NS      b.edu-servers.net.
```

```
edu.          172800    IN      NS      c.edu-servers.net.
```

```
...
```

```
;; ADDITIONAL SECTION:
```

```
a.edu-servers.net. 172800    IN      A       192.5.6.30
```

```
b.edu-servers.net. 172800    IN      A       192.33.14.30
```

```
c.edu-servers.net. 172800    IN      A       192.26.92.30
```

```
...
```

pick any one (multiple for redundancy)

this NS record says that a.edu-servers.net is a .edu name server.

# DNS lookup walkthrough

```
$ dig +norecurse eecs.berkeley.edu @198.41.0.4
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 26114
```

```
;; flags: qr; QUERY: 1, ANSWER: 0, AUTHORITY: 13, ADDITIONAL: 27
```

```
;; QUESTION SECTION:
```

```
;eecs.berkeley.edu.          IN      A
```

```
;; AUTHORITY SECTION:
```

```
edu.          172800    IN      NS      a.edu-servers.net.
```

```
edu.          172800    IN      NS      b.edu-servers.net.
```

```
edu.          172800    IN      NS      c.edu-servers.net.
```

```
...
```

```
;; ADDITIONAL SECTION:
```

```
a.edu-servers.net. 172800    IN      A      192.5.6.30
```

```
b.edu-servers.net. 172800    IN      A      192.33.14.30
```

```
c.edu-servers.net. 172800    IN      A      192.26.92.30
```

```
...
```

next nameserver IP to  
contact

# DNS lookup walkthrough

```
$ dig +norecurse eecs.berkeley.edu @192.5.6.30
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 36257
```

```
;; flags: qr; QUERY: 1, ANSWER: 0, AUTHORITY: 3, ADDITIONAL: 5
```

```
;; QUESTION SECTION:
```

```
;eecs.berkeley.edu.          IN      A
```

```
;; AUTHORITY SECTION:
```

```
berkeley.edu.      172800    IN      NS      adns1.berkeley.edu.
```

```
berkeley.edu.      172800    IN      NS      adns2.berkeley.edu.
```

```
berkeley.edu.      172800    IN      NS      adns3.berkeley.edu.
```

```
;; ADDITIONAL SECTION:
```

```
adns1.berkeley.edu. 172800    IN      A      128.32.136.3
```

```
adns2.berkeley.edu. 172800    IN      A      128.32.136.14
```

```
adns3.berkeley.edu. 172800    IN      A      192.107.102.142
```

```
...
```

# DNS lookup walkthrough

```
$ dig +norecurse eecs.berkeley.edu @192.5.6.30
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 36257
```

```
;; flags: qr; QUERY: 1, ANSWER: 0, AUTHORITY: 3, ADDITIONAL: 5
```

```
;; QUESTION SECTION:
```

```
;eecs.berkeley.edu.          IN      A
```

```
;; AUTHORITY SECTION:
```

```
berkeley.edu.               172800  IN      NS      adns1.berkeley.edu.
```

```
berkeley.edu.               172800  IN      NS      adns2.berkeley.edu.
```

```
berkeley.edu.               172800  IN      NS      adns3.berkeley.edu.
```

```
;; ADDITIONAL SECTION:
```

```
adns1.berkeley.edu. 172800  IN      A      128.32.136.3  
adns2.berkeley.edu. 172800  IN      A      128.32.136.14  
adns3.berkeley.edu. 172800  IN      A      192.107.102.142
```

```
...
```

next nameserver IP to  
contact

# DNS lookup walkthrough

```
$ dig +norecurse eecs.berkeley.edu @128.32.136.3
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 52788
```

```
;; flags: qr aa; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1
```

```
;; QUESTION SECTION:
```

```
;eecs.berkeley.edu.          IN      A
```

```
;; ANSWER SECTION:
```

```
eecs.berkeley.edu.  86400    IN      A      23.185.0.1
```

# DNS lookup walkthrough

```
$ dig +norecurse eecs.berkeley.edu @128.32.136.3
```

```
;; Got answer:
```

```
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 52788
```


```
;; flags: qr aa; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1
```

```
;; QUESTION SECTION:
```

```
eecs.berkeley.edu.          IN      A
```

```
;; ANSWER SECTION:
```

```
eecs.berkeley.edu. 86400    IN      A      23.185.0.1
```



one A type record in answer  
section: `eecs.berkeley.edu`'s  
IP is `23.185.0.1`.



# attack: cache poisoning

```
$ dig +norecurse eecs.berkeley.edu @128.32.136.3

;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 52788
;; flags: qr aa; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; QUESTION SECTION:
;eecs.berkeley.edu.          IN      A

;; ANSWER SECTION:
eecs.berkeley.edu.  86400   IN      A      23.185.0.1

;; ADDITIONAL SECTION:
www.google.com.      172800  IN      A      6.6.6.6
```

We made a query to a malicious berkeley.edu name server...

...and it returned a malicious record for www.google.com!

# attack: cache poisoning

```
$ dig +norecurse eecs.berkeley.edu @128.32.136.3

;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 52788
;; flags: qr aa; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; QUESTION SECTION:
;eecs.berkeley.edu.          IN      A

;; ANSWER SECTION:
eecs.berkeley.edu.  86400   IN      A      23.185.0.1

;; ADDITIONAL SECTION:
www.google.com.      172800  IN      A      6.6.6.6
```

We made a query to a malicious berkeley.edu name server...

...and it returned a malicious record for www.google.com!

now google.com has IP 6.6.6.6 in our cache! :(

# defense: bailiwick checking

- **bailiwick checking:** the resolver only accepts records if they are in the name server's zone

# defense: bailiwick checking

- **bailiwick checking:** the resolver only accepts records if they are in the name server's zone
  - ex: berkeley.edu NS can provide a record for eecs.berkeley.edu, but not mit.edu

# defense: bailiwick checking

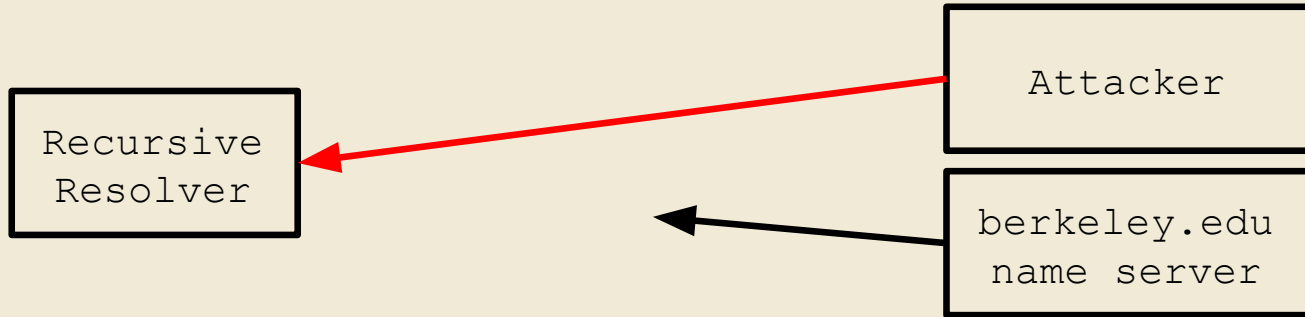
- **bailiwick checking:** the resolver only accepts records if they are in the name server's zone
  - ex: berkeley.edu NS can provide a record for eecs.berkeley.edu, but not mit.edu
- ex: root name server can provide a record for any domain (everything is in bailiwick for the root)

# attack: MITM, on-path

- MITM: can change DNS answer, poison cache

# attack: MITM, on-path

- MITM: can change DNS answer, poison cache
- on-path: can spoof DNS response, no fields to guess



# attack: MITM, on-path

- MITM: can change DNS answer, poison cache
- on-path: can spoof DNS response, no fields to guess
- off-path: need to guess 16-bit ID number



# Kaminsky attack

- the attacker includes
  - ``
  - ``
  - ``
  - ``
- client makes request for domain name for each one
- off-path attacker can guess ID every time & poison

# defenses (off-path)

- source port randomization (have to guess 32 bits)
- random domain query capitalization
- glue validation (don't cache glue records)
  - not implemented by all DNS software

**worksheet**  
(on 161 website)

# DNSSEC

duh, no security sucks. err, cryptography?

# securing DNS

# securing DNS

- DNS over TLS?

# securing DNS

- DNS over TLS?
  - slow

# securing DNS

- DNS over TLS?
  - slow
  - doesn't defend against malicious nameservers



# securing DNS

- DNS over TLS?
  - slow
  - doesn't defend against malicious nameservers
- we want integrity—no cache poisoning, tampering

# securing DNS

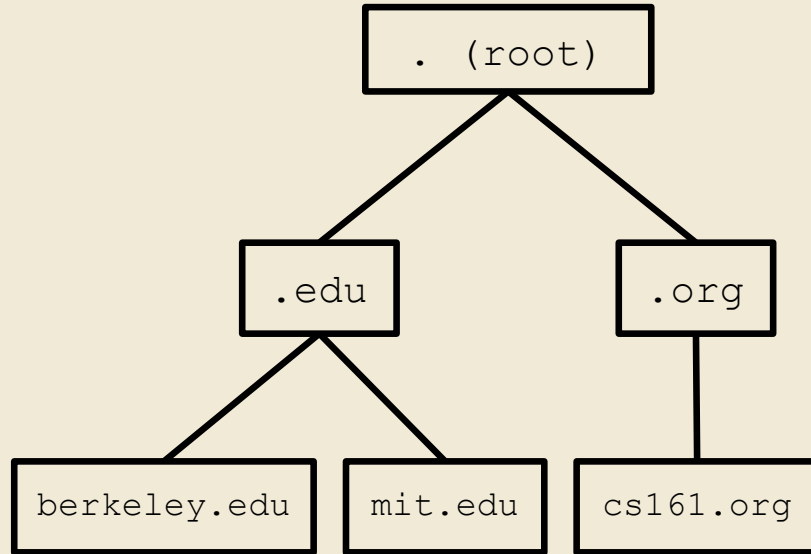
- DNS over TLS?
  - slow
  - doesn't defend against malicious nameservers
- we want integrity—no cache poisoning, tampering
- we don't need confidentiality
  - why?

# securing DNS

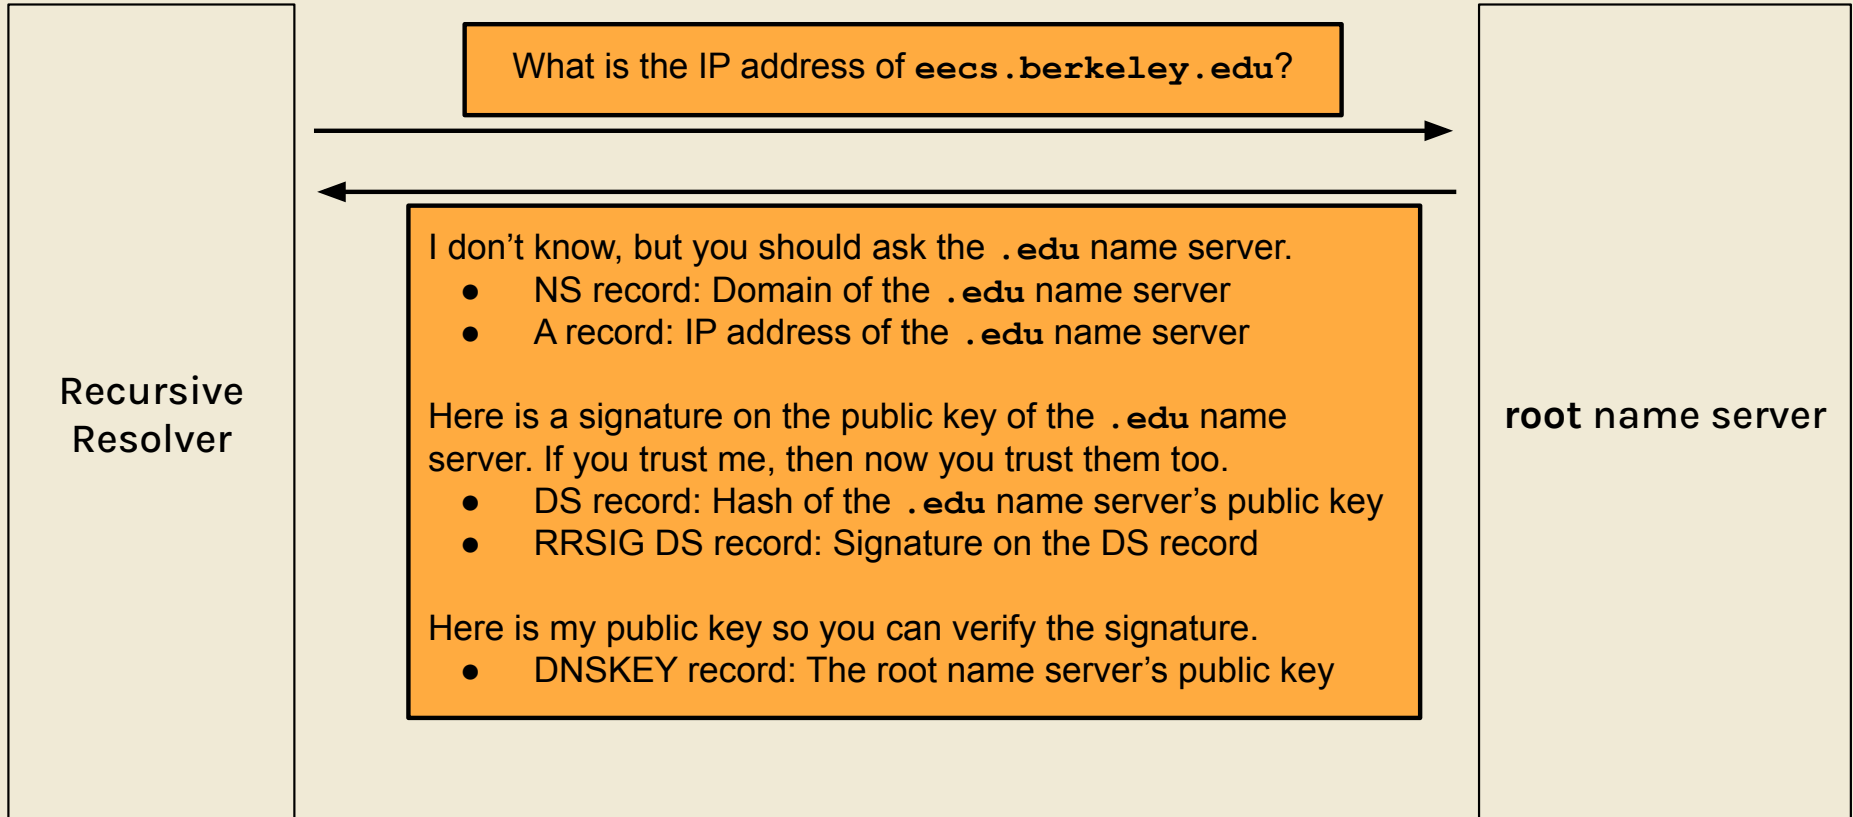
- DNS over TLS?
  - slow
  - doesn't defend against malicious nameservers
- we want integrity—no cache poisoning, tampering
- we don't need confidentiality
  - why?
  - anyone can make a DNS lookup, attackers can see IPs you connect to anyway

# DNSSEC (DNS Security Extensions)

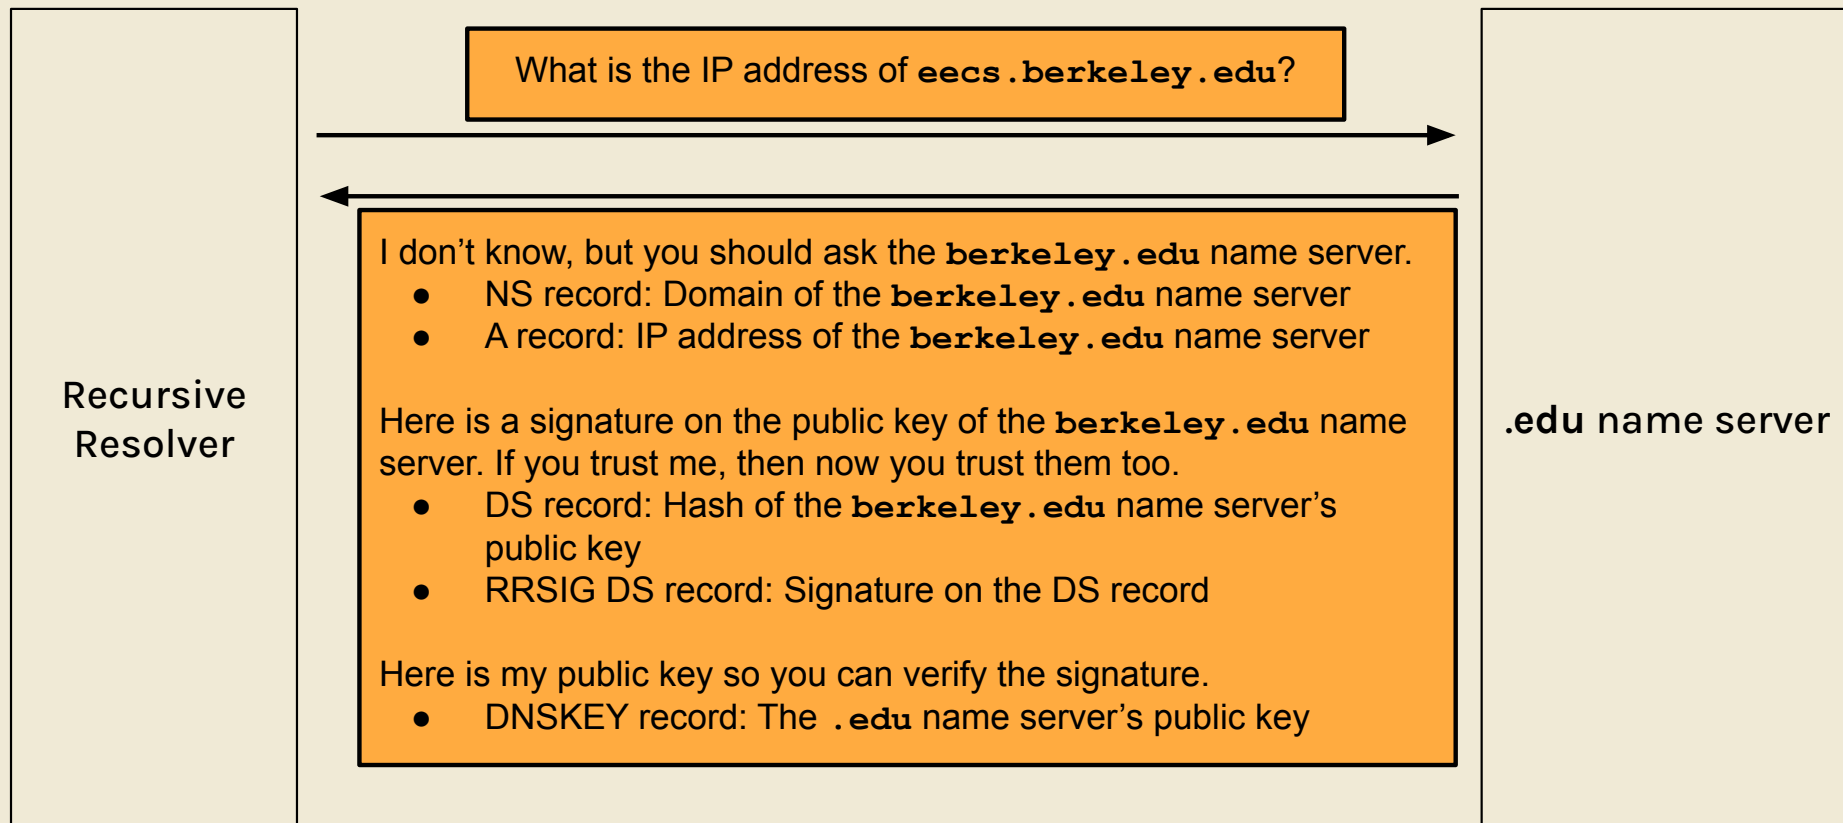
- delegate trust—parent signs child's public key



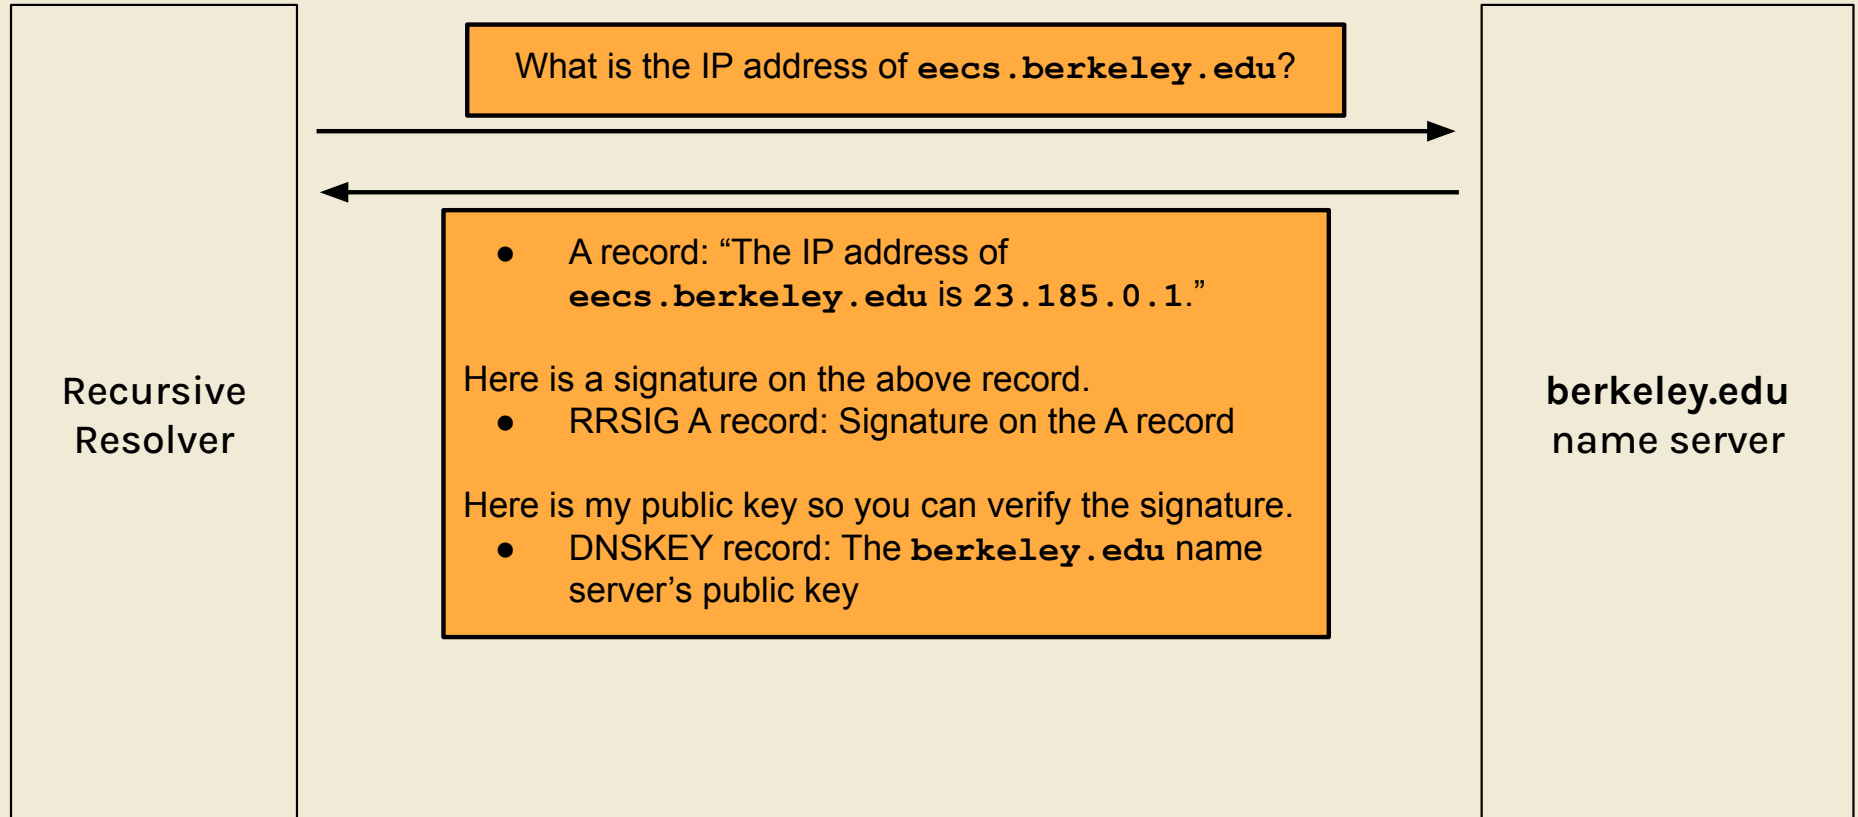
# DNSSEC lookup



# DNSSEC lookup



# DNSSEC lookup



# nonexistent domains

- what do we do if no records exist for query?



# nonexistent domains

- what do we do if no records exist for query?
- do we sign nothing and send it over?

# nonexistent domains

- what do we do if no records exist for query?
- do we sign nothing and send it over?
  - option 1: don't authenticate NXDOMAIN
    - bad (anyone can spoof)

# nonexistent domains

- what do we do if no records exist for query?
- do we sign nothing and send it over?
  - option 1: don't authenticate NXDOMAIN
    - bad (anyone can spoof)
  - option 2: sign a NXDOMAIN response
    - signing in real time is slow

# NSEC

- query for nonexistent.google.com

valid domains



maps

one

web

# NSEC

- query for nonexistent.google.com
- response: “no domains exist between maps.google.com and one.google.com”
  - can pre-sign all pairs of adjacent domains

valid domains



---

maps


one

web

# NSEC problem

- domain enumeration: can get every pair of domains by asking for arbitrary domains

valid domains



---

maps

one

web

# NSEC3

# NSEC3

- hash domain names and store pairs of adjacent hashes (instead of adjacent domain names)



# NSEC3

- hash domain names and store pairs of adjacent hashes (instead of adjacent domain names)

valid (hashed)  
domains



---

c612f3

d810de

# NSEC3

- hash domain names and store pairs of adjacent hashes (instead of adjacent domain names)
- query for nonexistent.google.com hashes to d48678...

valid (hashed)  
domains



---

c612f3

d810de

# NSEC3

- hash domain names and store pairs of adjacent hashes (instead of adjacent domain names)
- query for nonexistent.google.com hashes to d48678...
  - NSEC3 response: “There exist no domains which hash to values between c612f3... and d810de...”

valid (hashed)  
domains



---

c612f3

d810de

# hack of the day

- DNS resolver insecurities found on thousands of websites out of 7000 sampled
  - small/big businesses, governmental services
  - “...at least 25 were using static source ports... thousands of more domains using static source ports were discovered.”
  - none of the 25 used DNSSEC, etc.
  - cache poisoning used to hijack WordPress



feedback

[bit.ly/abhifedback](https://bit.ly/abhifedback)

slides: [bit.ly/cs161-disc](https://bit.ly/cs161-disc)